

MPhil Thesis

Michael Jing Long Lee

May 25, 2025

Contents

1	Introduction	1
1.1	Contributions	3
2	Background	4
2.1	Metaprogramming	4
2.1.1	Metaprogramming for Fast and Maintainable Code	4
2.1.2	The Design Space of Metalanguages	8
2.2	Effect Handlers	9
2.2.1	Composable and Customisable Effects	9
2.2.2	λ_{op} : A Calculus for Effect Handlers	11
2.2.3	The Design Space of Effect Handlers	19
2.3	Scope Extrusion	21
2.3.1	Existing Solutions to the Scope Extrusion Problem	22
3	Calculus	27
3.1	The Source Language: $\lambda_{\langle\langle\text{op}\rangle\rangle}$	27
3.1.1	Type System	28
3.2	The Core Language: $\lambda_{\text{AST}(\text{op})}$	35
3.2.1	Operational Semantics	36
3.2.2	Type System	39
3.2.3	Implementation	40
3.3	Elaboration from $\lambda_{\langle\langle\text{op}\rangle\rangle}$ to $\lambda_{\text{AST}(\text{op})}$	40
3.3.1	Elaborating Types	41
3.3.2	Elaborating Contexts	42
3.3.3	Elaborating Terms	42
3.3.4	Elaborating Typing Judgements	42
3.4	Metatheory	43
4	Scope Extrusion	46
4.1	Lazy Dynamic Check	46
4.2	Eager Dynamic Check	47
4.2.1	Correctness of the Eager Dynamic Check	48
4.2.2	Expressiveness of the Eager Dynamic Check	50
4.2.3	Efficiency of the Eager Dynamic Check	51
4.3	Best-Effort Dynamic Check	51
4.3.1	Correctness of the Best-Effort Dynamic Check	54
4.3.2	Expressiveness of the Best-Effort Dynamic Check	54
4.3.3	Efficiency of the Best-Effort Dynamic Check	55
4.4	Refined Environment Classifiers	55
4.4.1	Correctness of Refined Environment Classifiers	58
4.4.2	Expressiveness of Refined Environment Classifiers	62
4.5	Evaluation of $\lambda_{\langle\langle\text{op}\rangle\rangle}$	62

5	Conclusion	65
5.1	Limitations and Future Work	65

Abstract

Metaprogramming and effect handlers interact in unexpected, and sometimes undesirable, ways. One example is scope extrusion: the generation of ill-scoped code. There are many different ways to solve the problem of scope extrusion, but to my knowledge, until now there has been no way to evaluate them against each other. This dissertation introduces such a mechanism, which I use to evaluate the correctness, expressiveness, and efficiency of existing scope extrusion solutions. Additionally, I introduce a novel solution, a dynamic Best-Effort check, that I show is correct, and occupies a goldilocks zone between expressiveness and efficiency.

1 Introduction

Many compilers perform a wide range of optimisations. This empowers programmers to focus on writing maintainable code, entrusting the responsibility of *optimising* said code to the compiler.

While compiler optimisation is often all one needs, it is difficult for compiler engineers to ensure, *a priori*, that every user of the language has all the optimisations they require for their specific use case [24, 25]. When users discover that their desired optimisation has not been implemented, what can they do, beyond submitting a pull request?

Metaprogramming (for example, C++ templates) [1] allows the user to “teach the compiler a class of [new] tricks” [28], and is therefore one possible solution. In languages that support metaprogramming, programmers can identify what should be executed at compile-time. Importantly, programmers may write **maintainable** code, which when executed by the compiler, generates **efficient** code. For example, assume two implementations of the same function, one that runs faster on Arm, and one on Intel.

The maintainable way to switch between these functions is via a conditional:

```
1 let f () = if (cpu_type == Arm) then
2             (* return this program *)
3             else
4             (* return this program *)
5 f ()
```

OCaml

However, it might be expensive to determine processor type at run-time. If efficiency is paramount, programmers might have to duplicate their code, doubling maintenance costs. With metaprogramming, however, we may write

```
1 (* macro and $ can be read as annotations which
2    indicate that the code should be run at compile-time *)
3 macro m () = if (cpu_type == Arm) then
4               (* generate this program, to be run later *)
5               else
6               (* gen this program, to be run later *)
7 $(m ())
```

Macrocaml

Importantly, the macro `m` is executed by the compiler, *at compile time*. The Arm and Intel variants are *generated from a single specification*, not manually duplicated.

To support metaprogramming, languages have to provide the ability to build programs *to be run later*. These suspended programs are best thought of as abstract syntax trees. The following code constructs the program $(\lambda x.x)1$.

```

1  if (cpu_type == Arm) then
2    let build_app (f: (int -> int) expr) (v: int expr) = App(f, v)
3    let id_ast: (int -> int) expr = Lam(Var(x), Var(x))
4    build_app(id_ast, Int(1)) (* int expr *)
5  else
6    (*...*)

```

OCaml

Effect handlers are a powerful language construct that can simulate many other language features (state, I/O, greenthreading) [23], and have recently been added to OCaml [29]. It is thus strategic, and timely, to investigate the interaction between metaprogramming and effect handlers.

Metaprogramming is known to interact unexpectedly with effects. This interaction can be extremely desirable, for example, allowing programmers perform loop-invariant code-motion during code generation (on top of existing optimisations) [14] without making major modifications to their existing code [19].

Unfortunately, the interaction can also be undesirable. One problem is *scope extrusion* [14]. Scope extrusion occurs when the programmer accidentally generates code with unbound variables. In the following program, effects and metaprogramming combine to extrude `Var(x)` beyond its scope. The program stores `Var(x)` in a heap cell (1), and subsequently retrieves it outside its scope. The program thus evaluates to `Var(x)`, which is unbound.

```

1  let l: int expr ref = new(Int(1)) in
2  let id_ast : (int -> int) expr = Lam(Var(x), l := Var(x); Int(1)) in
3  !l

```

OCaml

The problem of scope extrusion has been widely studied, resulting in multitudinous mechanisms for managing the interaction between metaprogramming and effects. Some solutions adapt the type system (Refined Environment Classifiers [17, 10], Closed Types [5]), others insert dynamic checks into the generated code [14]. However, there are two issues.

First, I am not studying scope extrusion in a vacuum, but rather am working towards a concrete objective. The MacoCaml [34] project aims to extend the OCaml programming language, which has effect handlers, with compile-time metaprogramming. However, until now, the MacoCaml project has had no clear policy on how metaprogramming and effects should interact. In this setting, the options become more limited. Type-based approaches require unfeasible modification of the OCaml type-checker, and tend to limit expressiveness, disallowing a wide range of programs that do not lead to scope extrusion. Dynamic solutions are inefficient, reporting scope extrusion long after the error occurs, or unpredictable, allowing some programs, but disallowing morally equivalent programs. Further, many of the dynamic solutions have not been proved correct.

Second, and more importantly, I lacked a way to evaluate solutions fairly and holistically. Our evaluation criterion is three-pronged: correctness, efficiency, and expressiveness. Each solution is described in its own calculus, which made solutions difficult to compare. It is non-trivial to show that the definitions of scope extrusion in differing calculi agree. Hence, one solution may be correct with respect to its own definition, but wrong with respect to another. Further, expressiveness and efficiency are inherently comparative criteria.

This thesis solves both these problems. In [Chapter 3](#), I design a common language for

encoding and evaluating different solutions. In [Chapter 4](#), I use the designed language to formally describe and evaluate various checks. This process led to a novel Best-Effort dynamic check ([Section 4.3](#)), which I will argue occupies a goldilocks zone, and should be adopted by MacoCaml.

1.1 Contributions

Concretely, the contributions of this work are:

1. A novel two-stage calculus, $\lambda_{\langle\langle\text{op}\rangle\rangle}$, that allows for effect handlers at both stages: compile-time and run-time ([Chapter 3](#)). $\lambda_{\langle\langle\text{op}\rangle\rangle}$ is well-typed, has the expected metatheoretic properties ([Section 3.4](#)), and is designed to facilitate comparative evaluation of different scope extrusion checks ([Chapter 4](#)).
2. In $\lambda_{\langle\langle\text{op}\rangle\rangle}$, a formal description of the MetaOCaml check described by Kiselyov [14] ([Section 2.3.1](#)), as well as an evaluation of its correctness ([Section 4.2.1](#)), expressiveness ([Section 4.2.2](#)), and efficiency ([Section 4.2.3](#)).
3. In $\lambda_{\langle\langle\text{op}\rangle\rangle}$, a formal description of a novel Dynamic Best-Effort check ([Section 4.3](#)), which I prove correct ([Section 4.3.1](#)), and argue, by comparison with other checks, occupies a goldilocks zone between expressiveness ([Section 4.3.2](#)) and efficiency ([Section 4.3.3](#)).
4. An encoding of refined environment classifiers [17] into $\lambda_{\langle\langle\text{op}\rangle\rangle}$, with a proof of correctness via logical relation ([Section 4.4.1](#)), and an evaluation of its expressiveness compared to other checks ([Section 4.4.2](#)).
5. Implementations of the three dynamic checks in MacoCaml.

2 Background

The aim of this dissertation is to evaluate, and propose, policies for mediating the interaction between MacoCaml-style **metaprogramming** and **effect handlers**, by addressing the issue of **scope extrusion**.

In this chapter, I provide technical overviews of each of the key concepts: metaprogramming ([Section 2.1](#)), effect handlers ([Section 2.2](#)), and scope extrusion ([Section 2.3](#)).

2.1 Metaprogramming

What is MacoCaml-style metaprogramming? I will provide an answer in two steps. First, I motivate metaprogramming, by illustrating the challenge of writing code that is both fast and maintainable ([Section 2.1.1](#)). Second, I will consider the design space of metaprogramming ([Section 2.1.2](#)), highlighting decisions made by MacoCaml.

2.1.1 Metaprogramming for Fast and Maintainable Code

Metaprogramming helps programmers write fast and maintainable code. How does one write fast and maintainable code? A naïve answer is “by being a skilled programmer”¹. Programmer skill is insufficient, because maintainability and efficiency are in constant tension.

I illustrate this tension by considering a concrete problem. Consider computing the gradient of a differentiable function as part of backpropagation over a neural network. More precisely, assume f of type differentiable

```
1 type differentiable = Sin | Tanh | Sigmoid | ...
2                   | Polynomial of float list
3                   | Compose of differentiable * differentiable
```

OCaml

For example, the following expression represents $\sin \circ \tanh$.

```
1 Compose(Tanh, Sin)
```

OCaml

We wish to write a function `grad` such that $\text{grad } f = f'$. For simplicity, assume the existence of a helper function, `grad_of`, that returns the gradient of basic functions. For example, `grad_of Sin 0.0 = cos 0.0 = 1.0`. The `grad_main` function ([Listing 1](#)) is one maintainable way to compute gradients:

¹Though not the worst answer: “use ChatGPT”


```

1  let rec grad_main f x = match f with
2    | Sin
3    | Tanh
4    | Sigmoid
5    | ...
6    | Polynomial(_) -> grad_of f x
7    | Compose(f, g) -> (grad_main f x) * (grad_main g (app f x))

```

OCaml

Listing 1: A maintainable implementation of grad

However, `grad_main` may not be the most efficient implementation. Performing a `match` on every recursive call might result in expensive branches. If x is a vector, and the weights of a polynomial are vectors, then `grad_main` could hide opportunities for cache prefetching.

If f is known in advance, for example, $f = \text{Compose}(\text{Tanh}, \text{Sin})$, we could implement a more efficient `grad_fast` function (Listing 2), whose body is simply a hardcoded equation:

```

1  let grad_fast x = (cos x) /. (cosh (sin x) ** 2)

```

OCaml

Listing 2: A fast implementation of grad, assuming $f = \text{Compose}(\text{Tanh}, \text{Sin})$

Although `grad_fast` only works for a single f , it has eliminated the branching overhead, and enabled opportunities for prefetching. It is thus likely to be faster.

The grad example illustrates the trade-off between maintainability and efficiency. `grad_main` is maintainable in part because it parameterises over f . More generally, abstraction centralises implementations, thus reducing maintenance costs. However, `grad_fast` is more efficient because it is more specialised to a specific f . More generally, many compiler optimisations, like monomorphisation, eliminate abstraction, simplifying functions by applying known arguments in advance.

The tension between maintainability (abstraction) and efficiency (specialisation) has also been observed in regex matching [31], parsing [37], linking [27], statistical modelling [33], and hardware design [32].

A more informed answer might therefore be “by letting the compiler generate optimised versions of my maintainable code”. Not quite: for reasons both theoretical and practical, compiler optimisations can be insufficient. In theory, we proposed an optimisation that assumed we would always know f at, or before, compile-time. Is this a reasonable assumption? It is: we assumed that grad would perform backpropagation over neural networks. The network over which backpropagation is performed is known at compile-time. However, notice that this justification appeals to domain-specific knowledge regarding how grad will be used. In the general case, grad could be applied to a function not known until runtime. It is not feasible to expect a compiler to spot all opportunities for optimisation [24]. In practice, while compiler engineers might have an economic incentive to write optimisations for the machine learning community, this may not be true for less lucrative domains [25]. Even in machine learning, many libraries are built on top of existing languages, like Python, which might not perform the desired optimisations.

How does one write maintainable and efficient code, **when one cannot trust the compiler to optimise one’s code?**

2 Background

One answer is metaprogramming, which gives users the ability to perform code-generation. Programmers may thus take matters into their own hands: manually generating optimised code when the compiler may not automatically do so for them. The `grad` function in JAX, a Python-based machine learning framework, uses metaprogramming for precisely this purpose [11].

Speeding up exponentiation with Metaprogramming

Metaprogramming allows for code-generation via the creation and manipulation of abstract syntax trees (ASTs). I will now illustrate how metaprogramming works with reference to `MacoCaml`, which implements *compile-time* metaprogramming.

While the `grad` example motivated metaprogramming, for pedagogical reasons, I switch to a morally equivalent, but simpler example: raising an integer x to an exponent n . One maintainable implementation is the `pow` function (Listing 3):

```
1 let rec pow (n: int) (x: int) =  
2   if n == 0 then 1  
3   else x * pow (n-1) x
```

OCaml

Listing 3: A maintainable implementation of an exponentiation function

As it may be applied to any exponent n , `pow` is analogous to `grad_main` (Listing 1), which could be applied to *any* differentiable function f .

However, should we know the exponent in advance, for example $n = 2$, then a more efficient, but less maintainable implementation, is the `square` function (Listing 4), analogous to `grad_fast` (Listing 2).

```
1 let square x = x * x
```

OCaml

Listing 4: An efficient implementation of exponentiation, assuming $n = 2$

Metaprogramming can be utilised to write a function, `pow_gen`, which resembles `pow` (inheriting its maintainability), but that generates a program which resembles `square` (inheriting its efficiency). Listing 5 presents the meta-programmed `pow_gen` function. Compilation generates `y * y * 1`, which resembles the body of `square`. I will now explain the mechanics of generation.

```
1 macro rec pow_gen (n: int) (x: int expr) =  
2   if n == 0 then <<1>>  
3   else <<$x * $(pow_gen (n-1) x)>>  
4   let square y = $(pow_gen 2 <<y>>) (*after compile-time: y * y * 1 *)  
5   square 3 (*at runtime: 9*)
```

OCaml

Listing 5: A meta-programmed `pow_gen` function, which resembles `pow` but generates `square`

Recall that (compile-time) metaprogramming gives the programmer the ability to generate programs at compile-time, for use at run-time. We may build this in two steps, by:

1. Deciding on a representation for code values, such that code can be created and manipulated by programs. Once we have a representation for code values, it is possible to write expressions that return code values. These expressions serve as program generators.
2. Building a mechanism for executing expressions *at compile-time*. We can constrain this mechanism, using types, so only generators can be executed at compile-time.

First, we represent code values as ASTs. Generated programs are ASTs, and program generators are expressions that evaluate to ASTs. For clarity, assume that there exists an AST node for each language construct. For example, the integer 1 has AST node `Int(1)`. If a program has type 'a, then its AST node has type 'a expr. One can now write program generators, that evaluate to ASTs, for example:

```
1 let rec pow_gen (n: int) (x: int expr) =
2   if n == 0 then Int(1)
3   else Mul(x, (pow_gen (n-1) x))
4 pow_gen 2 Int(3) (*Mul(Int(3), Mul(Int(3), 1))**)
5 pow_gen 2 Var(y) (*Mul(Var(y), Mul(Var(y), 1))**)
6 pow_gen 3 Var(y) (*Mul(Var(y), Mul(Var(y), Mul(Var(y), 1))))*)
```

Macrocaml

Second, we need a mechanism to execute expressions at compile-time. In Macrocaml, this is the “top-level splice”, a splice (\$) annotation not surrounded by quotes (<<>>). In Listing 5, there is only one top-level splice, on line 4: \$(pow_gen 2 <<y>>). We may now shift program generators (and only program generators) under top-level splices, to perform generation at compile time. Note that to access pow_gen at compile-time, we must also move it under the top-level splice.

```
1 let square y = $(let rec pow_gen (n: int) (x: int expr) = ...
2   in pow_gen 2 Var(y))
3   (*Mul(Var(y), Mul(Var(y), 1))**)
4 let cube y = $(let rec pow_gen (n: int) (x: int expr) = ...
5   in pow_gen 3 Var(y))
6   (*Mul(Var(y), Mul(Var(y), Mul(Var(y), 1))))*)
```

Macrocaml

To allow compile-time functions, like pow_gen, to be re-used across multiple top-level splices, Macrocaml introduces the **macro** (Listing 6)

```
1 macro rec pow_gen (n: int) (x: int expr) =
2   if n == 0 then Int(1)
3   else Mul(x, (pow_gen (n-1) x))
4 let square y = $(pow_gen 2 Var(y)) (*Mul(Var(y), Mul(Var(y), 1))**)
5 let cube y = $(pow_gen 3 Var(y)) (*Mul(Var(y), Mul(Var(y), Mul(Var(y), 1))))*)
```

OCaml

Listing 6: In Macrocaml, **macro** allows for definitions to be shared across top-level splices

Further, rather than explicit AST constructors, ASTs are created by the <<>> (“quote”) and \$ annotations. Quotation creates ASTs, by converting a program into its AST representation. For example,

```
<< $x + 0 >> = Plus(Var(x), Int(0))
```

2 Background

Under a quotation, the `$` annotation stops this conversion, allowing for programs that *manipulate* ASTs.

```
<< $x + 0 >> = Plus(x, Int(0))
```

In MacoCaml, the programmer interleaves quotes and splices to perform code generation

```
<< $(add_zero <<1>>) + 0 >> = Plus(add_zero Int(1), Int(0))
```

Notice that `$` is overloaded. We must be careful to disambiguate between “top-level splices”, which execute programs at compile-time, and splices under quotations, which stop conversion to AST.

Re-writing Listing 6 in this style (being careful about non-top-level splices), we obtain exactly Listing 5.

Applying this technique to the grad example, we obtain

```
1 macro rec grad_gen f x = match f with
2   | Sin
3   | Tanh
4   | Sigmoid
5   | ...
6   | Polynomial(_) -> grad_of f x
7   | Compose(f, g) -> <<$(grad_gen f x) * $(grad_gen g (app f x))>>
8 let grad_fast x = $(grad_gen Compose(Tanh, Sin) <<x>>)
```

MacoCaml

where `grad_of` and `app` are appropriately modified.

2.1.2 The Design Space of Metalanguages

Different metalanguages provide slightly different variants of metaprogramming to the user. In this section, I taxonomise these languages by considering three key design decisions:

1. Homogenous or Heterogenous

Do the generated (“object”) and generating (“meta”) languages agree or differ?

If the object and meta languages are the same, this is known as homogenous metaprogramming. Otherwise, it is heterogenous [15].

MacoCaml allows for homogenous metaprogramming, where OCaml code generates OCaml code. In contrast, MetaHaskell [20] programs generate C code, allowing for heterogenous metaprogramming.

2. Run-time or Compile-Time

When does the generation take place?

Code generation could take place at compile-time (as with MacoCaml programs or C macros), or at run-time (as with MetaOCaml [14]).

Run-time and compile-time metaprogramming differ non-trivially. The former requires a language construct (`!`, or “run”) for explicit invocation of the compiler. Further, with run-time metaprogramming, generated and generating programs may share a heap.

MacoCaml supports compile-time metaprogramming, and we will pay no further attention to run-time metaprogramming.

3. Two-stage or Multi-stage

How many stages of code generation are allowed?

When introducing MacoCaml, I illustrated how one uses top-level splices to shift computation from run-time (“level 0”) to compile-time (“level -1 ”). Might it be possible to shift computation from compile-time to a pre-compile-time (“level -2 ” phase), for example, via a nested splice?

```
1  $( $ pow_gen 2 Var(y) )
```

MacoCaml

In a two-stage system, one is restricted to operating between two levels, so this is disallowed. In contrast, in a multi-stage system, one can operate between any number of levels. Multi-stage metaprogramming is thus strictly more general than two-stage metaprogramming.

Although nested splices are disallowed in MacoCaml, it is a multi-stage system, since entire modules may be imported at a decremented level [35].

MacoCaml offers homogenous, compile-time, multi-stage metaprogramming. The scope of this dissertation is slightly more restrictive: I focus on two-stage, not multi-stage metaprogramming. This restriction was motivated by a cost-benefit analysis:

1. **Cost:** Since in MacoCaml, the module system is the only mechanism for achieving multi-stage programming, investigating multi-stage metaprogramming would require the investigation of module systems, effects, and metaprogramming. The interaction between module systems and metaprogramming is still an ongoing area of research [6].
2. **Benefit:** In practice, “almost all uses” of multi-stage metaprogramming only use two stages [9]. Further, scope extrusion can be observed, and is often studied, in two-stage systems [10, 17].

2.2 Effect Handlers

What is an effect handler? I will first motivate effect handlers by considering the problem of adding resumable exceptions to OCaml ([Section 2.2.1](#)). Second, I will introduce a calculus for studying the operational behaviour of effect handlers, à la Pretnar [23] ([Section 2.2.2](#)). This calculus will be useful both for precise description of effect handlers, and as a basis for investigating the interaction between metaprogramming and effect handlers (once the calculus has been extended with metaprogramming facilities). Finally, since different design decisions for effect handlers could affect the nature of their interaction with metaprogramming, I will consider the design space of effect handlers ([Section 2.2.3](#)).

2.2.1 Composable and Customisable Effects

Effects are a mechanism by which a program interacts with its environment. Examples of effects include state, (resumable) exceptions, non-determinism, and I/O. Effects are

2 Background

typically defined and understood separately, meaning they are not easily composable. They are also implemented by compiler engineers rather than programmers, meaning they are not customisable. Effect handlers provide a programmable, unifying framework that may be instantiated into different effects. This allows for composable and customisable treatment of effects.

To illustrate the need for effect handlers, consider the following problem, by Kiselyov [13]. Assume a binary search tree of (key, value) pairs. The following code provides two functions. The first finds a value v associated with key k , raising a `NotFound` exception if k is not in the tree. The second updates the dictionary with a fresh key value pair, overwriting old values.

```
1  type ('a, 'b) tree = Lf | Br of 'a * 'b * tree * tree
2
3  let rec find (t: tree) (k: 'a) = match t with
4  | Lf -> raise NotFound()
5  | Br(k', v, l, r) -> if k == k' then v
6                        else if k < k' then find l k
7                        else find r k
8
9  let rec update (t: tree) (k: 'a) (v: 'b) = match t with
10 | Lf -> Br(k, v, Lf, Lf)
11 | Br(k', v', l, r) -> if k == k' then Br(k, v, l, r)
12                       else if k < k' then Br(k', v', update l k v, r)
13                       else Br(k', v', l, update r k v)
```

OCaml

Assume the task is to build a `findOrInsert` function that either finds the value associated with a key, *or* inserts a default value. A naïve approach to writing this function would be

```
1  let rec findOrInsert (t: tree) (k: 'a) (default: 'b) =
2    try find t k with NotFound -> insert t k default
```

OCaml

This function is **inefficient**. If the `find` function raises a `NotFound` exception, it will do so at the point where the default value should be inserted. If computation could be resumed at the point where the exception was raised, as such:

```
1  let rec findOrInsert (t: tree) (k: 'a) (default: 'b) =
2    try find t k with NotFound(p) -> continue p Br(k, default, Lf, Lf)
```

OCaml

then the function could be twice as fast. p represents the suspended program to be resumed, and is known as a *delimited continuation*.

The aforementioned problem motivates the need for resumable exceptions. To understand the need for effect handlers, consider how one might go about **implementing** resumable exceptions. One approach might be to fork the implementation of handlers and tweak it ever-so-slightly. This solution does not scale well. First, the solution may not be **composable**. The intended informal semantics for resumable exceptions is “effectively equivalent to exceptions, with the additional power to resume programs”. Resumable exceptions should thus interact with other exceptions in a predictable way, but this is difficult to guarantee, and *continually* guarantee, especially as implementations evolve, and more variants of exceptions are demanded. Second, the solution is not **cus-**

tomisable. To add resumable exceptions requires a compiler engineer to modify the compiler. With the exception of raising an issue, there is nothing the programmer may do, in the moment, to meet their need.

Effect handlers resolve both composability and customisability issues. Much like how exception handlers allow users to create custom exceptions with custom semantics, effect handlers provide a general framework for creating custom effects with custom semantics. The interaction between effect handlers is described abstractly, parameterising over the exact semantics of the effect. Hence, implementing effects (in the earlier example, resumable exceptions, but more generally, state, I/O, greenthreading, non-determinism, and more) as effect handlers ensures composability by design.

With effect handlers, we can re-write the previous example to obtain the behaviour of resumable exceptions, even if the OCaml compiler does not support it, with the guarantee that **NotFound** will interact predictably with other defined effects.

```

1  type _ Effect.t += NotFound: unit -> tree t
2
3  let rec find (t: tree) (k: 'a) = match t with
4    | Lf -> NotFound()
5    | Br(k', v, l, r) -> if k == k' then v
6                        else if k < k' then find l k
7                        else find r k
8
9  let rec findOrInsert (t: tree) (k: 'a) (default: 'b) =
10    match find t k with NotFound(p) with
11    | v -> v
12    | effect NotFound k -> continue p Br(k, default, Lf, Lf)

```

OCaml

Since effect handlers may be instantiated into a range of different effects, considering the interaction of metaprogramming with effect handlers is an exercise in killing many birds with a single stone. Additionally, effect handlers were recently added to OCaml [29], making their interaction a timely problem.

2.2.2 λ_{op} : A Calculus for Effect Handlers

Having motivated effect handlers, I will now describe a calculus, which I call λ_{op} , for reasoning about their operational behaviour. λ_{op} is a slight variant of the calculus described by Pretnar [23]. Understanding λ_{op} will be useful for two reasons. First, it will aid reasoning about the interaction between effects and metaprogramming. Second, my universal calculus will be described by extending λ_{op} .

```

handle
  do  $x \leftarrow \text{print}(1)$ ; return 1 in do  $y \leftarrow \text{print}(2)$ ; return 2 in  $x + y$ 
with
  {return( $x$ )  $\mapsto$  return ( $x$ , "");
   print( $x, k$ )  $\mapsto$  do ( $v, s$ )  $\leftarrow$  continue  $k$  () in return ( $v, f"\{x\};" ^ s$ )}

return (3, "1;2")

```

 λ_{op}

Listing 7: An λ_{op} program that returns (3, "1;2"). It will be used as a running example throughout this section.

Syntax λ_{op}

Values	$v ::= x \mid n \mid \lambda x.c \mid \kappa x.c$
Computations	$c ::= v_1 v_2 \mid \mathbf{return} \ v \mid \mathbf{do} \ x \leftarrow c_1 \ \mathbf{in} \ c_2$ $\quad \mid \mathbf{op}(v) \mid \mathbf{handle} \ c \ \mathbf{with} \ \{h\} \mid \mathbf{continue} \ v_1 \ v_2$
Handlers	$h ::= \mathbf{return}(x) \mapsto c \mid h; \mathbf{op}(x, k) \mapsto c$

Figure 2.1: The syntax of λ_{op} . Terms are syntactically divided into values v , computations c , and handlers h

Figure 2.1 collates the base syntax of λ_{op} . In addition to this base syntax, in this section, I will assume λ_{op} is extended with the following language extensions: a unit value $()$, pairs $(1, 2)$ which can be destructured $\mathbf{do} \ (x, y) \leftarrow \mathbf{return} \ (1, 2) \ \mathbf{in} \ x + y$, strings "Hello", format strings $f\{1\}$, and string concatenation \wedge . For example, the following code evaluates to "Revolution 9".

```
do (x, y) ← return ("Revolution", f"{9}") in x ^ y
return "Revolution 9"
```

 λ_{op}

Further, I use $c_1; c_2$ as syntactic sugar for $\mathbf{do} \ _ \leftarrow c_1 \ \mathbf{in} \ c_2$. I will explain key language constructs in turn, with reference to a running example: the λ_{op} program in Listing 7.

Sequencing computations: do and return

Effects force us to carefully consider the order of evaluation. For example, consider the following OCaml programs

```
1 let pure      = (1+0) + (2+0)
2 let effectful = let l = new 0 in (l := 1; 1) + (l := 2; 2)
```

OCaml

The result of `pure`, which has no effects, is independent of the evaluation order. In contrast, the result of `effectful` is dependent on the evaluation order. If terms are evaluated left-to-right, the value of `!l` is 2, otherwise, it is 1.

In order to be precise about the order of evaluation, λ_{op} terms are stratified into two syntactic categories, "inert values" v and "potentially effectful computations" c [23]. $\mathbf{return} \ v$ lifts values into computations, and is also the result of fully evaluating a computation. $\mathbf{do} \ x \leftarrow c_1 \ \mathbf{in} \ c_2$ sequences computations, forcing programmers to be explicit the order of evaluation. First, c_1 is fully evaluated to obtain some $\mathbf{return} \ v$. The value v is then bound to x , and finally c_2 is evaluated.

For example, extending λ_{op} with a plus function, what is the order of evaluation of plus 1 2? Are both arguments evaluated before application, or are evaluation and application interleaved? The syntax forces programmers to choose explicitly. The programmer can either fully evaluate both arguments before applying them in turn,

```
do x ← return 1 in (do y ← return 2 in (do f ← plus x in fy))
```

 λ_{op}

or alternatively, evaluate 1, apply it, then evaluate 2


```
do x ← return 1 in (do f ← plus x in do y ← return 2 in fy)
```



Both choices are valid, but the programmer must choose. For clarity, where the ordering cannot affect the result (both of the aforementioned choices evaluate to **return** 3), I will abuse notation and write (for instance) $1 + 2$.

Performing effects: **op**, **handle**, and **continue**

Having made explicit the order of operation, I now consider effect handlers. Recall that effect handlers allow users to register custom effects with custom semantics. I will now illustrate how this is supported by λ_{op} .

For simplicity, λ_{op} assumes that the effects have been registered in advanced, parameterising over them with the placeholder **op**(v). Assume that the user has declared the effects **print** and **read_int** in advance. This would allow the user to write programs like

```
do x ← print(1); return 1 in do y ← print(2); return 2 in x + y
```



In the program fragment above, **print** is an effect, but with as-yet-unknown semantics. Effect handlers, which comprise a **return handler** and zero or more **operation handlers**, specify how effects interact with their environment, and thus may be used to give effects meaning. I will define an effect handler that accumulates print statements in a string (some “stdout”). For example, the aforementioned program should return (3, “1; 2”).

I begin by considering how to handle the case where there are no calls to **print**. For example, in the program **return** 3. I wish to return both the value, and the empty string (empty stdout) to the environment: in this case, (3, “”). I achieve this by specifying a *return handler*.

$$\text{return}(x) \mapsto c$$

In this case, I set c to **return** (x , “”). All effect handlers must specify a return handler. In many cases, the return handler is simply the identity (c is set to **return** x).

Next, I consider how to handle a call to **print**. I use an operation handler of the form

$$\text{print}(x, k) \mapsto c$$

Where c is the user-defined semantics for **print**. Concretely, one instance of c is

$$\text{print}(x, k) \mapsto \text{do } (v, s) \leftarrow \text{continue } k () \text{ in return } (v, f''\{x\}; " ^ s)$$

In the definition of c , the programmer may refer to x and k , which I will now explain. x allows programs to send values (for example, values to be printed) to their environment. k is a delimited continuation representing a suspended program, awaiting a value from the environment. Effects also allow programs to receive data from their environment, as in

$$1 + \text{get_int_from_user}()$$

Note that the program is suspended until the value is received. I write the suspended program as $1 + [-]$, where $[-]$ indicates an as-yet-unknown value. This suspended program is represented by the continuation k . The expression

$$\text{continue } k v$$

is used to resume the suspended program with value v .

I now able to interpret the concrete operation handler c : the handler resumes the suspended program, supplying a unit value, since **print** effects do not receive values from their environment. This returns a value v and some partially accumulated stdout s . The handler then prepends the printed value, x , onto s .

Having defined the semantics for **print**, the user may now interpret the earlier example with their semantics, using the **handle** e **with** $\{h\}$ construct. Doing so results in the program in [Listing 7](#).

Notice that multiple effects may be handled by the same handler, and the same effect might be handled by multiple handlers, potentially with different semantics.

Operational Semantics

Having described λ_{op} informally, I now describe its semantics more precisely, via an operational semantics [Figure 2.2](#).

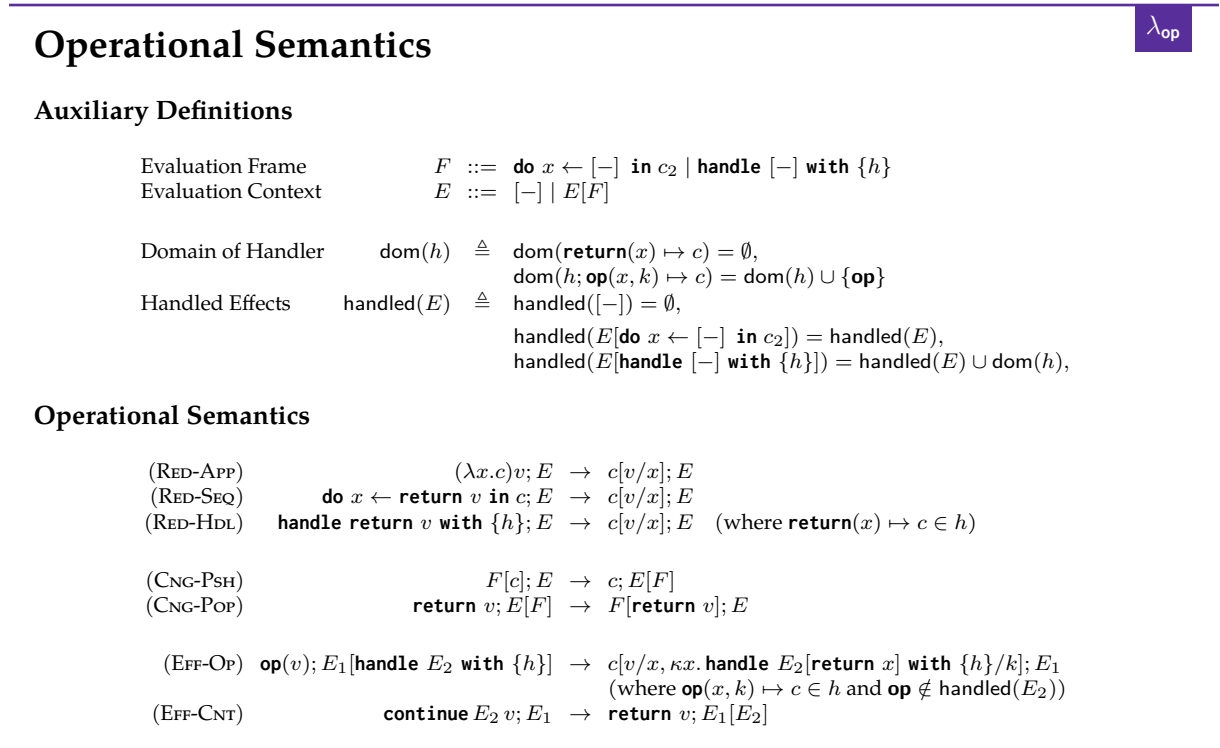


Figure 2.2: The operational semantics of λ_{op} . The semantics is given on configurations of the form $\langle c, E \rangle$, with the brackets dropped for clarity. Rules are divided into three classes: reduction rules RED- X , which perform computation, congruence rules CNG- Y which manipulate the evaluation context, and effect rules EFF- Z that are special to λ_{op}

The operational semantics is given on configurations of the form $\langle c, E \rangle$, where c is a term and E is an evaluation context, in the style of Felleisen and Friedman [7]. Evaluation contexts are represented as a stack of evaluation frames F , à la Kiselyov [13]. Most of the rules are standard. I focus on two rules: EFF-OP, the mechanism for giving effects custom semantics, and EFF-CNT, the mechanism for resuming programs.

To illustrate the operation of **EFF-OP** and **EFF-CNT**, consider the evaluation of the running example in [Listing 7](#), beginning with an empty context. Let h be the handler body

```
{return( $x$ )  $\mapsto$  return ( $x$ , "");
 print( $x$ ,  $k$ )  $\mapsto$  do ( $v$ ,  $s$ )  $\leftarrow$  continue  $k$  () in return ( $v$ , f"{ $x$ };" ^  $s$ )}
```

After several applications of **CNG-Psh**, we obtain the configuration

```
< print(1) ; handle
    do  $x \leftarrow [-]$ ; return 1 in do  $y \leftarrow$  print(2); return 2 in  $x + y$ 
  with { $h$ } >
```

Let $E = \text{do } x \leftarrow \text{return } u; \text{return } 1 \text{ in } \text{do } y \leftarrow \text{print}(2); \text{return } 2 \text{ in } x + y$. Applying **EFF-OP**, we can suspend the program, find the handler h with the user's semantics for **print**, and give the **print** effect the desired semantics

```
< do ( $v$ ,  $s$ )  $\leftarrow$  continue ( $\kappa u$ . handle  $E$  with { $h$ }) () in return ( $v$ , f"{1};" ^  $s$ ) ; [-] >
```

Applying **CNG-Psh**,

```
< continue ( $\kappa u$ . handle  $E$  with { $h$ }) () ; do ( $v$ ,  $s$ )  $\leftarrow [-]$  in return ( $v$ , f"{1};" ^  $s$ ) >
```

Applying **EFF-CNT**, we can resume the program that was suspended

```
< return () ; do( $v$ ,  $s$ )  $\leftarrow$ 
    handle
        do  $x \leftarrow$  ([-]; return 1) in
        do  $y \leftarrow$  (print(2); return 2)
        in  $x + y$ 
    with { $h$ }
  in return ( $v$ , f"{1};" ^  $s$ ) >
```

The side-condition on **EFF-OP** is needed because the user may define multiple handlers with different semantics for the same effect. The side-condition resolves any ambiguity by using the *latest* handler. For example, the following program has a **read** effect that is given two definitions: it could read either 1 or 2. The ambiguity is resolved by choosing the latest handler: in this case, 1.

```
handle
  handle read() with {return( $y$ )  $\mapsto$  return  $y$ ; read( $x$ ,  $k$ )  $\mapsto$  continue  $k$  1}
with
  {return( $y$ )  $\mapsto$  return  $y$ ; read( $x$ ,  $k$ )  $\mapsto$  continue  $k$  2}

return 1
```

λ_{op}

Type-and-Effect System

I now give a type-and-effect system to λ_{op} . [Figure 2.3](#) collates the syntax of λ_{op} types, which I will now briefly describe.

I will assume a universe of pre-declared operations, Σ , where for any effect op , Σ can be used to look up the type of op .

$$\text{op} : A \rightarrow B \in \Sigma$$

Types

Effects row	$\Delta ::= \emptyset \mid \Delta \cup \{\text{op}_i\}$
Value type	$T ::= \mathbb{N}$ $\mid T_1 \xrightarrow{\Delta} T_2$ functions $\mid T_1 \xrightarrow{\Delta} T_2$ continuations
Computation type	$T! \Delta$
Handler type	$T_1! \Delta \Longrightarrow T_2! \Delta'$

Figure 2.3: λ_{op} types. Notice that, just as terms are divided into values, computations, and handlers, types are divided into value types (T), computation types ($T! \Delta$), and handler types ($T_1! \Delta \Longrightarrow T_2! \Delta'$)

I will not impose any restrictions on A and B . In particular, they may be recursive, and refer to op . For example:

$$\text{recursive} : 1 \rightarrow (1 \xrightarrow{\{\text{recursive}\}} 1)$$

Thus, one may write programs which do not terminate, by “tying the knot”.

Just like terms, types are divided into value types (for example, \mathbb{N}), computation types ($\mathbb{N}! \{\text{print}\}$), and handler types ($\mathbb{N}! \{\text{print}\} \Longrightarrow \mathbb{N}! \emptyset$). Since computations may have effects, computation types track unhandled effects using an effects row (Δ), which in this system is simply a set. This type-and-effect system allows us to distinguish between values, computations that return values, and computations that return values and additionally have some unhandled side effects.

Term	Type
3	\mathbb{N}
do $x \leftarrow \text{return } 1$ in do $y \leftarrow \text{return } 2$ in $x + y$	$\mathbb{N}! \emptyset$
do $x \leftarrow \text{print}(1); \text{return } 1$ in do $y \leftarrow \text{return } 2$ in $x + y$	$\mathbb{N}! \{\text{print}\}$

Functions are values, and are applied to other values, but produce computations on application. For example, the function

$$\lambda x : \mathbb{N}. \text{print}(x); \text{return } x$$

is a value that accepts a value of type \mathbb{N} and returns a computation of type $\mathbb{N}! \{\text{print}\}$. Functions thus have suspended effects, which I write $T_1 \xrightarrow{\Delta} T_2$. In this case, the function has type $\mathbb{N} \xrightarrow{\{\text{print}\}} \mathbb{N}$. For technical reasons, continuations and functions need to be distinguished, but in most cases they may be treated equivalently.

Handlers transform computations of one type to computations of another type. This happens in two ways: first, by handling effects, and thus removing them from the effects row (which recall represents unhandled effects). Second, by modifying the return type of computations. To reflect both abilities, handlers are given a type of the form $T_1! \Delta \Longrightarrow T_2! \Delta'$. For example, a handler of the form

$$\begin{aligned} &\{\text{return}(x) \mapsto \text{return } (x, ""); \\ &\quad \text{print}(x, k) \mapsto \text{do } (v, s) \leftarrow \text{continue } k() \text{ in return } (v, f''\{x\}; " ^ s)\} \end{aligned}$$

Typing Rules

(NAT)	(VAR)	(LAMBDA)	(CONTINUATION)
$\frac{}{\Gamma \vdash n : \mathbb{N}}$	$\frac{\Gamma(x) = T}{\Gamma \vdash x : T}$	$\frac{\Gamma, x : T_1 \vdash c : T_2 ! \Delta}{\Gamma \vdash \lambda x.c : T_1 \multimap T_2}$	$\frac{\Gamma, x : T_1 \vdash c : T_2 ! \Delta}{\Gamma \vdash \kappa x.c : T_1 \multimap T_2}$
(APP)	(CONTINUE)		
$\frac{\Gamma \vdash v_1 : T_1 \multimap T_2 \quad \Gamma \vdash v_2 : T_1}{\Gamma \vdash v_1 v_2 : T_2 ! \Delta}$	$\frac{\Gamma \vdash v_1 : T_1 \multimap T_2 \quad \Gamma \vdash v_2 : T_1}{\Gamma \vdash \mathbf{continue} v_1 v_2 : T_2 ! \Delta}$		
(RETURN)	(DO)		
$\frac{\Gamma \vdash v : T}{\Gamma \vdash \mathbf{return} v : T ! \Delta}$	$\frac{\Gamma \vdash c_1 : T_1 ! \Delta \quad \Gamma, x : T_1 \vdash c_2 : T_2 ! \Delta}{\Gamma \vdash \mathbf{do} x \leftarrow c_1 \mathbf{in} c_2 : T_2 ! \Delta}$		
(OP)	(HANDLE)		
$\frac{\Gamma \vdash v : T_1 \quad \mathbf{op} : T_1 \rightarrow T_2 \in \Sigma \quad \mathbf{op} \in \Delta}{\Gamma \vdash \mathbf{op}(v) : T_2 ! \Delta}$	$\frac{\Gamma \vdash c : T_1 ! \Delta \quad \Gamma \vdash h : T_1 ! \Delta \Longrightarrow T_2 ! \Delta' \quad \forall \mathbf{op} \in \Delta \setminus \Delta'. \mathbf{op} \in \text{dom}(h)}{\Gamma \vdash \mathbf{handle} c \mathbf{with} \{h\} : T_2 ! \Delta'}$		
(RET-HANDLER)			
$\frac{\Gamma, x : T_1 \vdash c : T_2 ! \Delta'}{\Gamma \vdash \mathbf{return}(x) \mapsto c : T_1 ! \Delta \Longrightarrow T_2 ! \Delta'}$			
(OP-HANDLER)			
$\frac{\mathbf{op} : A \rightarrow B \in \Sigma \quad \Gamma \vdash h : T_1 ! \Delta \Longrightarrow T_2 ! \Delta' \quad \Gamma, x : A, k : B \xrightarrow{\Delta'} T_2 \vdash c : T_2 ! \Delta' \quad \Delta' \subseteq \Delta \setminus \{\mathbf{op}\} \quad \mathbf{op}(x', k') \mapsto c' \notin h}{\Gamma \vdash h; \mathbf{op}(x, k) \mapsto c : T_1 ! \Delta \Longrightarrow T_2 ! \Delta'}$			

Figure 2.4: Typing rules for λ_{op} terms

may be given type $\mathbb{N} ! \{\text{print}\} \Rightarrow (\mathbb{N} \times \text{String}) ! \emptyset$, reflecting both the handling of the **print** effect and the transformation of the return type to include the collated print statements.

I now consider the typing rules for terms, which are collated in Figure 2.4. Most rules are standard, but a few are worth paying attention to.

First, the RETURN and DO rules. In the RETURN rule, we are allowed to assign the term **return** v any set of effects. For example, we could write:

$$\frac{}{\Gamma \vdash \text{return } \emptyset : \mathbb{N} ! \{\text{print}\}}$$

This flexibility is important, because to type **do** $x \leftarrow c_1$ **in** c_2 , the DO rule requires both c_1 and c_2 to have the same effects. For example, without this flexibility, we would not be able to complete the following typing derivation

$$\frac{\vdots}{\Gamma \vdash \text{do } x \leftarrow \text{print}(\emptyset) \text{ in } \text{return } \emptyset : \mathbb{N} ! \{\text{print}\}}$$

2 Background

A valid alternative would be to forbid this flexibility and add explicit subtyping. However, such an approach would no longer be syntax directed.

Second, the **OP** rule. Previously, I assumed that the user declared their effects in advance. I additionally assume that they declare the types of their effects in advance, and that I store the mapping from effects to types in Σ . For example, I might assume $\Sigma = \{\mathbf{print} : \mathbb{N} \rightarrow 1\}$. In OCaml, this would correspond to writing:

```
1  type _ Effect.t += Print: nat -> unit OCaml
```

Note further the $\text{op} \in \Delta$ restriction – flexibility allows us to over-approximate the effects in a term, but never underapproximate them.

Third, the **RET-HANDLER** and **OP-HANDLER** rules, which are used to type handlers, which I will explain by means of an example. Assume we are trying to type the handler

$$\begin{aligned} \{\mathbf{return}(x) \mapsto \mathbf{return}(x, ""); \\ \mathbf{print}(x, k) \mapsto \mathbf{do}(v, s) \leftarrow \mathbf{continue} k() \text{ in } \mathbf{return}(v, f''\{x\}; " \wedge s)\} \end{aligned}$$

with the type $\mathbb{N}! \{\mathbf{print}\} \Longrightarrow (\mathbb{N} \times \text{String})! \emptyset$. We apply the **OP-HANDLER** rule, which is transcribed below. Preconditions are numbered for reference.

$$\begin{array}{c} (\text{OP-HANDLER}) \\ \frac{\begin{array}{l} (1) \text{op} : A \rightarrow B \in \Sigma \\ (2) \Gamma \vdash h : T_1! \Delta \Longrightarrow T_2! \Delta' \\ (3) \Gamma, x : A, k : B \xrightarrow{\Delta'} T_2 \vdash c : T_2! \Delta' \quad (4) \Delta' \subseteq \Delta \setminus \{\text{op}\} \quad (5) \text{op}(x', k') \mapsto c' \notin h \end{array}}{\Gamma \vdash h; \text{op}(x, k) \mapsto c : T_1! \Delta \Longrightarrow T_2! \Delta'} \end{array}$$

The preconditions of the **OP-HANDLER** rule direct us to check, in turn:

- (1) $\mathbf{print} : \mathbb{N} \rightarrow 1 \in \Sigma$, which is true by assumption
- (2) Recursively check the rest of the handler $h = \mathbf{return}(x) \mapsto \mathbf{return}(x, "")$, ensuring it has type $\mathbb{N}! \{\mathbf{print}\} \Longrightarrow (\mathbb{N} \times \text{String})! \emptyset$. This follows from a trivial application of the **RET-HANDLER** rule.
- (3) Assuming x has type \mathbb{N} and k has type $1 \xrightarrow{\emptyset} (\mathbb{N} \times \text{String})$, the body

$$\mathbf{do}(v, s) \leftarrow \mathbf{continue} k() \text{ in } \mathbf{return}(v, f''\{x\}; " \wedge s)$$

has type $(\mathbb{N} \times \text{String})! \emptyset$. This is easy to show.

- (4) That the handler *only* removes **print** from the effects row, and no other effects. This check passes, but would fail if we tried to type the handler with, for example, $\mathbb{N}! \{\mathbf{print}, \mathbf{get}\} \Longrightarrow (\mathbb{N} \times \text{String})! \emptyset$
- (5) That there are no other handlers for **print** in h .

Finally, I define a notion of well-typed computation.

Definition 2.2.1 (Well-typed computation) c is well-typed if $\cdot \vdash c : T! \emptyset$

Metatheory

I will build not only on λ_{op} , but on metatheoretic properties of λ_{op} , which are proven by Bauer and Pretnar [2].

Theorem 2.2.1 (Progress) *If $\Gamma \vdash E[c] : T! \Delta$ then either*

1. *c of the form **return** v and $E = [-]$,*
2. *c of the form **op**(v) for some $\text{op} \in \Delta$,*
3. *$\exists.E', c'$ such that $\langle c; E \rangle \rightarrow \langle c'; E' \rangle$*

Theorem 2.2.2 (Preservation) *If $\Gamma \vdash E[c] : T! \Delta$ and $\langle c; E \rangle \rightarrow \langle c'; E' \rangle$, then $\Gamma \vdash E'[c'] : T! \Delta$*

Corollary 2.2.1 (Type Safety) *If $\cdot \vdash c : T! \emptyset$ then either*

1. *$\langle c; [-] \rangle \rightarrow^\omega$ (non-termination)*
2. *$\langle c; [-] \rangle \rightarrow^* \langle \text{return } v; [-] \rangle$*

2.2.3 The Design Space of Effect Handlers

The design space of effect handlers is large. I consider three key design decisions made by different systems.

1. Named or Unnamed Handlers

Can I invoke a specific handler for an operation?

In λ_{op} , when there are multiple handlers for the same effect, the “nearest” or “most recent” handler for that effect is invoked. An alternative approach is **named handlers** [34], where, by associating each handler with a *name*, the programmer can more easily specify which handler should be invoked.

While named handlers can be more ergonomic, they do not provide greater expressiveness than using distinct effects [34]. I do not consider named handlers in this thesis.

2. Deep, Shallow, or Sheep Handlers

Are multiple instances of the same effect handled by the same handler?

In λ_{op} , continuations reinstate handlers (EFF-CNT) and thus multiple instances of the same effect are handled by the same handler. For example, in the following example, the effect **addn** is handled by the same handler, adding one each time. These are known as **deep** handlers.

```

handle
  addn(1) + addn(2)
with
  {return( $x$ )  $\mapsto$  return  $x$ ;
   addn( $y, k$ )  $\mapsto$  continue  $k$  ( $y + 1$ )}
```

return 5

 λ_{op}

I could also have chosen *not* to reinstate the handler, in an approach known as **shallow** handlers [8]. The example above would be stuck, since the second **addn** would not be handled.

Finally, I could have chosen to modify the interface for **continue** such that it accepts a handler

$$\mathbf{continue} \ k \ v \ h$$

This would allow multiple effects to be handled by different handlers. That is, the programmer could add 1 the first time **addn** is performed, and 2 the second time. These handlers behave as a hybrid of shallow and deep handlers, and are thus termed **sheep** handlers [21].

OCaml allows the programmer to choose between shallow and deep handlers. Since most prior work on scope extrusion focuses on deep handlers [10], I focus on those.

3. One-Shot or Multi-Shot Continuations

How many times can one resume the same continuation?

In λ_{op} , continuations may be resumed multiple times. For example, the programmer can write

```

handle
  performTwice(1)
with
  {return( $x$ )  $\mapsto$  return  $x$ ;
   performTwice( $y, k$ )  $\mapsto$  (continue  $k$   $y$ ) + (continue  $k$   $y$ )}
```

return 2

 λ_{op}

I say the effect system permits **multi-shot continuations**. Multi-shot continuations are useful for simulating certain effects, like non-determinism [21].

In other systems, like OCaml and WasmFX [21], this is not allowed: continuations are only allowed to be resumed once. These systems permit **one-shot continuations**.

Although continuations in OCaml are one-shot, due to the utility of multi-shot continuations, I believe it is worthwhile to study effect systems with multi-shot continuations.

2.3 Scope Extrusion

I now turn my attention to scope extrusion, which arises from the unexpected interaction of effects and metaprogramming. To illustrate scope extrusion, I will first extend λ_{op} (Page 11) with AST constructors $\text{Var}(x_T)$, $\text{Nat}(n)$, Lam , and Plus . For example, we may generate the AST of $\lambda x : \mathbb{N}. x + 0$ as follows:

```
return Lam(Var( $x_{\mathbb{N}}$ ), Plus(Var( $x_{\mathbb{N}}$ ), Nat(0)))
```

λ_{op}

Listing 8 illustrates the problem of scope extrusion. The program constructs the AST of $\lambda x : \mathbb{N}. x$, but additionally performs an effect, **extrude**, with type $\mathbb{N} \text{ expr} \rightarrow \mathbb{N} \text{ expr}$. The handler for **extrude** discards the continuation, simply returning the value it was given: $\text{Var}(x_{\mathbb{N}})$. The entire program evaluates to $\text{Var}(x_{\mathbb{N}})$, and the generated AST is ill-scoped. We say that the result of evaluation demonstrates scope extrusion.

```
handle
  do body  $\leftarrow$  extrude(Var( $x_{\mathbb{N}}$ )) in return Lam(Var( $x_{\mathbb{N}}$ ), body)
with
  {return( $u$ )  $\mapsto$  return Nat(0);
   extrude( $y, k$ )  $\mapsto$  return  $y$ }

return Var( $x_{\mathbb{N}}$ )
```

λ_{op}

Listing 8: A λ_{op} program that evaluates to the $\text{Var}(x_{\mathbb{N}})$. The AST is ill-scoped, and thus exhibits scope extrusion. It will be used as a running example.

It is difficult to give a precise definition to scope extrusion, because there are multiple competing definitions [14, 17], and many are given informally. For example, is scope extrusion a property of the *result* of evaluation [17], as in Listing 8, or is it a property of *intermediate* configurations [14]? We can, for example, build ASTs with extruded variables, that are bound at some future point. In Listing 9, we produce the intermediate AST $\text{Plus}(\text{Nat}(0), \text{Var}(x_{\mathbb{N}}))$, which is not well scoped. However, the result of evaluation is well scoped: $\text{Lam}(\text{Var}(x_{\mathbb{N}}), \text{Plus}(\text{Nat}(0), \text{Var}(x_{\mathbb{N}})))$. Does Listing 9 exhibit scope extrusion?

```
handle
  do body  $\leftarrow$  extrude(Var( $x_{\mathbb{N}}$ )); return Var( $x_{\mathbb{N}}$ ) in return Lam(Var( $x_{\mathbb{N}}$ ), body)
with
  {return( $u$ )  $\mapsto$  return  $u$ ;
   extrude( $y, k$ )  $\mapsto$  do  $z \leftarrow$  return Plus(Nat(0),  $y$ ) in continue  $k$   $z$ }

return Lam(Var( $x_{\mathbb{N}}$ ), Plus(Nat(0), Var( $x_{\mathbb{N}}$ )))
```

λ_{op}

Listing 9: A λ_{op} program that may, or may not demonstrate scope extrusion, depending on one's definition. The final result of the program is well-scoped, but not all intermediate results are well-scoped.

Nevertheless, all definitions agree on the example in Listing 8. Making precise the competing definitions of scope extrusion these competing definitions, and their relation

to one another, is a contribution of this dissertation.

2.3.1 Existing Solutions to the Scope Extrusion Problem

There are multiple solutions to the problem of scope extrusion. The solution space can be broadly divided into two types of approaches: static (type-based) and dynamic. I will now survey two dynamic approaches, which I term the lazy and eager checks, and one static approach, the method of refined environment classifiers [17, 10].

Lazy Dynamic Check

Scope extrusion, at least, of the kind in Listing 12, may seem trivial to resolve: evaluate the program to completion, and check that the resulting AST is well-scoped [14]. I term this the Lazy Dynamic Check. This approach, while clearly correct and maximally expressive, is not ideal for efficiency and error reporting reasons.

To illustrate the inefficiency of this approach, consider a slight variation of Listing 8, Listing 10. In Listing 10, we can, in theory, report a warning as soon scope extrusion is detected. However, waiting for the result of the program can be much more inefficient.

In terms of error reporting, note that, in waiting for the result of execution, the lazy dynamic check loses information about *which program fragment* was responsible for scope extrusion, reducing the informativeness of reported errors [14].

```

do  $x \leftarrow$  handle
  do body  $\leftarrow$  extrude(Var( $x_{\mathbb{N}}$ )) in return Lam(Var( $x_{\mathbb{N}}$ ), body)
  with
    {return( $u$ )  $\mapsto$  return Nat(0);
     extrude( $y, k$ )  $\mapsto$  return  $y$ }
  in some very long program; return  $x$ 

return Var( $x_{\mathbb{N}}$ )

```

λ_{op}

Listing 10: A λ_{op} program that evaluates to the $\text{Var}(x_{\mathbb{N}})$. Executing the entire program to determine if it exhibits scope extrusion is inefficient.

Eager Dynamic Check

A second dynamic check, motivated by the problems with the Lazy Dynamic Check, adopts a stricter definition of scope extrusion. During the code generation process, one inserts checks into the running program, reporting errors when one encounters unbound free variables in intermediate ASTs. Hence, the Eager Dynamic Check would classify the program in Listing 9 as exhibiting scope extrusion. The Eager Dynamic Check has been adopted by BER MetaOCaml, and offers better efficiency and error reporting guarantees over the Lazy Dynamic Check [14].

However, the Eager Dynamic Check is not without issue. The problem relates to the manner in which checks are inserted, which I will make precise later. To illustrate the problem, consider Listing 11, a slight variation of Listing 9 in which the program fragment $\text{Plus}(\text{Nat}(0), y)$ is replaced with y .

```

handle
  do body  $\leftarrow$  extrude(Var( $x_{\mathbb{N}}$ )); return Var( $x_{\mathbb{N}}$ ) in return Lam(Var( $x_{\mathbb{N}}$ ), body)
with
  {return( $u$ )  $\mapsto$  return  $u$ ;
   extrude( $y, k$ )  $\mapsto$  do  $z \leftarrow$  return  $y$  in continue  $k\ z$ }

return Lam(Var( $x_{\mathbb{N}}$ ), Plus(Nat(0), Var( $x_{\mathbb{N}}$ )))

```

Listing 11: A λ_{op} program that is a slight variation of Listing 9, but that (unlike Listing 9) passes the Eager Dynamic Check.

While the program in Listing 11 produces an intermediate AST with extruded variables, because of the mechanism for inserting checks, it passes the Eager Dynamic Check. I assert that this behaviour is unintuitive, and exposes too much of the internals to the programmer.

Refined Environment Classifiers

Refined Environment Classifiers are a static check that uses the type system to prevent scope extrusion. Recall that metaprogramming involves the *creation* and *manipulation* of ASTs. Refined environment classifiers prevent scope extrusion by checking:

1. *Created* ASTs are well-scoped
2. *Manipulating* ASTs preserves well-scopedness

I shall first ignore the ability to manipulate ASTs, and consider how to ensure *created* ASTs are well-scoped. What does it mean to be well-scoped? Consider Figure 2.5, the AST of

$$(\lambda f. \lambda x. f x)(\lambda y. y)$$

Informally, a scope represents a set of variables that are permitted to be free. In the example, there are four scopes: one where no variables are free, one where only f is free, one where f and x are free, and one where y is free. An AST is well-scoped *at a scope* if it is well-typed, and where the only free variables are those permitted by the scope.

Refined environment classifiers make this notion precise. Each classifier represents a scope. The AST has four scopes, corresponding to four classifiers:

- γ_{\perp} The top-level, where no free variables are permitted
- γ_f Only Var(f) permitted to be free
- γ_{fx} Only Var(f) and Var(x) permitted to be free
- γ_y Only Var(y) permitted to be free

As every variable binder creates a scope, we may refer to the classifier created by Var(α), classifier(Var(α)). For example, $\gamma_{fx} = \text{classifier}(\text{Var}(x))$.

With classifiers, I may now be precise about “the variables permitted to be free (within the scope)”. As illustrated by the nesting in Figure 2.5, scopes are related to other




Figure 2.5: The AST of $(\lambda f.\lambda x.fx)(\lambda y.y)$, where each scope is labelled with the corresponding environment classifier.

scopes. For example, since the scope γ_{fx} is created within scope γ_f , any variable tagged with γ_f may be safely used in the scope γ_{fx} . I say γ_f is compatible with γ_{fx} , and write

$$\gamma_f \sqsubseteq \gamma_{fx}$$

The compatibility relation (\sqsubseteq) is a partial order, meaning \sqsubseteq is reflexive, anti-symmetric, and transitive, and identifies a smallest classifier. Reflexivity expresses that $\text{Var}(\alpha)$ may be used within the scope it creates

$$\forall \gamma. \gamma \sqsubseteq \gamma$$

anti-symmetric, since nesting only proceeds in one direction

$$\forall \gamma_1, \gamma_2. \gamma_1 \sqsubseteq \gamma_2 \wedge \gamma_2 \sqsubseteq \gamma_1 \implies \gamma_1 = \gamma_2$$

and transitive, since we should count nestings within nestings

$$\forall \gamma_1, \gamma_2, \gamma_3. \gamma_1 \sqsubseteq \gamma_2 \wedge \gamma_2 \sqsubseteq \gamma_3 \implies \gamma_1 \sqsubseteq \gamma_3$$

γ_\perp acts as the least element of this partial order

$$\forall \gamma. \gamma_\perp \sqsubseteq \gamma$$

Given a classifier γ , I may now define the variables permitted to be free in γ , written $\text{permitted}(\gamma)$

$$\text{permitted}(\gamma) \triangleq \{\text{Var}(\alpha) \mid \text{classifier}(\text{Var}(\alpha)) \sqsubseteq \gamma\}$$

For example, $\text{permitted}(\gamma_{fx}) = \{\text{Var}(f), \text{Var}(x)\}$

I say that an AST n is well-scoped at type T and scope γ if it is well-typed at T , and all free variables in n are in $\text{permitted}(\gamma)$. I write

$$\Gamma \vdash^\gamma n : T$$




Figure 2.6: Visual depiction of the (C-Abs) typing rule.

Ensuring that created ASTs are well-scoped is the responsibility of the type system. The key rule is the C-Abs rule, which, **assuming we know that we are creating an AST**, has roughly the following shape²:

$$\begin{array}{c}
 \text{(C-Abs)} \\
 \frac{\gamma \in \Gamma \quad (2) \gamma' \text{fresh} \quad (3) \Gamma, \gamma', \gamma \sqsubseteq \gamma', (x : T_1)^{\gamma'} \vdash^{\gamma'} n : T_2}{(1) \Gamma \vdash^{\gamma} \lambda x. n : T_1 \longrightarrow T_2}
 \end{array}$$

The premises and conclusions have been numbered for reference, and many technical details have been simplified for clarity. Figure 2.6 visually depicts the typing rule.

- (1) The goal of the typing rule is to ensure the function is well-scoped at type $T_1 \rightarrow T_2$ and scope γ .
- (2) Since the function introduces a new variable binder, $\text{Var}(x)$, one has to create a new scope. This is achieved by picking a fresh classifier γ' .
- (3) We record the following:
 - (a) Since γ' is created within the scope of γ , $\gamma \sqsubseteq \gamma'$.
 - (b) $\text{classifier}(\text{Var}(x)) = \gamma'$ (as a shorthand $(x : T_1)^{\gamma'}$)

With this added knowledge, we ensure that the function body is well-scoped at type T_2 and γ' .

The above example focused on **creating** ASTs, and had no compile-time executable code. I now consider how to maintain well-scopedness while **manipulating** ASTs. I consider the “AST” of the scope extrusion example, Listing 8 (Figure 2.7). Notice that in place of AST nodes, we may now have compile-time executable code that *evaluate* to AST nodes. Thus, both code and AST nodes reside within scopes. We have two classifiers: γ_{\perp} and γ_{α} , with $\text{classifier}(\text{Var}(\alpha)) = \gamma_{\alpha}$.

Key to the prevention of scope extrusion is the typing of handlers and operations, like **extrude**, that manipulate ASTs. Since these rules are complex, I describe them informally. The handle expression

handle e **with** $\{h\}$

is in scope γ_{\perp} . Therefore, for each operation handled by h , such as **extrude**, the argument to the operation must either not be an AST, or be an AST that is well-scoped at some $\gamma \sqsubseteq \gamma_{\perp}$. However, $\text{Var}(\alpha)$ is typed at γ_{α} , and clearly, $\gamma_{\alpha} \not\sqsubseteq \gamma_{\perp}$. There is thus no way to type the scope extrusion example in Listing 8.

²I will revisit this assumption when introducing the type system for my calculus




Figure 2.7: The “AST” of the scope extrusion example, Listing 8. Notice that in place of AST nodes, we may now have compile-time executable code that *evaluate* to AST nodes.

Note that the analysis was independent of the *body* of the handler. Therefore, the examples in Listings 9 and 11 are *also* not well-typed. Perhaps somewhat surprisingly, so too is Listing 12 (which would pass both the Eager and Lazy Dynamic Checks). Refined environment classifiers statically prevent variables ($\text{Var}(\alpha)$) from becoming *available* in program fragments ($\text{op}(y, k) \mapsto \dots$) where, *if misused, might* result in scope extrusion. This is, of course, an over-approximation. It means that the refined environment classifiers check prevents not only both types of scope extrusion, but even more benign examples, such as that in Listing 12.

```

handle
  do body  $\leftarrow$  extrude( $\text{Var}(x_{\mathbb{N}})$ ); return  $\text{Var}(x_{\mathbb{N}})$  in return  $\text{Lam}(\text{Var}(x_{\mathbb{N}}), \text{body})$ 
with
  { return( $u$ )  $\mapsto$  return  $\text{App}(u, \text{Nat}(1))$ ;
    extrude( $y, k$ )  $\mapsto$  return  $\text{Nat}(0)$  }

return  $\text{Nat}(0)$ 

```

λ_{op}

Listing 12: A λ_{op} program that passes the Eager and Lazy Dynamic Checks, but is not well-typed under the Refined Environment Classifiers type system.

Listing 12 thus illustrates one of the key drawbacks of Refined Environment Classifiers: the check is too stringent, and restricts expressiveness.

3 Calculus




Figure 3.1: $\lambda_{\langle\text{op}\rangle}$ is first elaborated into $\lambda_{\text{AST}(\text{op})}$, which is then executed **at compile-time** to obtain the AST of a run-time λ_{op} program.

To re-iterate, I am considering the interaction between homogenous, compile-time, two-stage **metaprogramming** (Page 4), and an **effect system** with deep handlers and multi-shot continuations (Page 9). In this chapter, I describe a calculus, $\lambda_{\langle\text{op}\rangle}$, for studying said interaction. $\lambda_{\langle\text{op}\rangle}$ will have both metaprogramming and effect handlers. To the best of my knowledge, this is the first calculus in which one may write effectful compile-time code that generates effectful run-time code. However, $\lambda_{\langle\text{op}\rangle}$ will not mediate the interaction between metaprogramming and effects: scope extrusion prevention is not a language feature. Rather, the aim will be to extend $\lambda_{\langle\text{op}\rangle}$ with various scope extrusion checks, such that the checks may be evaluated in a comparative fashion.

Programs written in $\lambda_{\langle\text{op}\rangle}$ cannot be directly executed. Rather, following the style of Xie et al. [35], one must first elaborate (or compile) from $\lambda_{\langle\text{op}\rangle}$ (the “source” language) to a “core” language, $\lambda_{\text{AST}(\text{op})}$. Programs written in $\lambda_{\text{AST}(\text{op})}$ may then be executed, to obtain the AST of a run-time λ_{op} program. This process is summarised in Figure 3.1. Elaboration is necessary, since it is the point at which dynamic checks may be inserted.

In this chapter, I will first introduce $\lambda_{\langle\text{op}\rangle}$ (Section 3.1), then $\lambda_{\text{AST}(\text{op})}$ (Section 3.2). Following this, I will describe the elaboration from $\lambda_{\langle\text{op}\rangle}$ to $\lambda_{\text{AST}(\text{op})}$ (Section 3.3). Finally, I will discuss the metatheoretic properties of $\lambda_{\langle\text{op}\rangle}$ (Section 3.4).

3.1 The Source Language: $\lambda_{\langle\text{op}\rangle}$

$\lambda_{\langle\text{op}\rangle}$ extends λ_{op} with quotes and splices. Recall that λ_{op} , following a fine-grain call-by-value approach, divides terms into two syntactic categories, values v and computations

Syntax		$\lambda_{\langle\text{op}\rangle}$
Values	$v ::= \dots$	
Expressions	$e ::= \dots \mid \langle\langle e \rangle\rangle \mid \e	

Figure 3.2: $\lambda_{\langle\text{op}\rangle}$ syntax. The syntax is broadly the same as λ_{op} , except with the addition of quotes and splices.

Effects Row

Run-Time $\xi ::= \cdot \mid \xi \cup \{\text{op}_i^0\}$
Compile-Time $\Delta ::= \cdot \mid \Delta \cup \{\text{op}_i^{-1}\}$

Types

Level 0	Values	$T^0 ::= \mathbb{N}^0$	naturals
		$\mid (T_1^0 \xrightarrow{\xi} T_2^0)^0$	functions
		$\mid (T_1^0 \xrightarrow{\xi} T_2^0)^0$	continuations
	Computations	$T^0 ! \xi$ $\mid T^0 ! \Delta; \xi$	
	Handlers	$(T_1^0 ! \xi \implies T_2^0 ! \xi') ! \Delta$	
Level -1	Values	$T^{-1} ::= \mathbb{N}^{-1}$	naturals
		$\mid (T_1^{-1} \xrightarrow{\Delta} T_2^{-1})^{-1}$	functions
		$\mid (T_1^{-1} \xrightarrow{\Delta} T_2^{-1})^{-1}$	continuations
		$\mid \text{Code}(T^0 ! \xi)^{-1}$	run-time code
	Computations	$T^{-1} ! \Delta$	
	Handlers	$T_1^{-1} ! \Delta \implies T_2^{-1} ! \Delta'$	

Figure 3.3: $\lambda_{\langle\text{op}\rangle}$ types. I highlight three important elements: first, types are stratified into two levels, 0 and -1. Second, effects are stratified into two levels, ξ (for run-time effects) and Δ for compile-time effects. Third, the Code type allows for compile-time programs to manipulate ASTs of run-time code.

c. $\lambda_{\langle\text{op}\rangle}$ is similar, dividing terms into values v and expressions e (Figure 3.7)

Notice that we cannot quote values: we *must* generate effectful programs. For example, $\langle\langle 1 \rangle\rangle$ is not valid syntax, instead, one must write $\langle\langle \text{return } 1 \rangle\rangle$. However, for clarity, I will abuse notation and write $\langle\langle 1 \rangle\rangle$. Similarly, $\langle\langle \text{return } 1 \rangle\rangle$ is a computation, not a value, so one must write

do $a \leftarrow \langle\langle \text{return } 1 \rangle\rangle$ **in** $\text{op}(a)$

rather than $\text{op}(\langle\langle \text{return } 1 \rangle\rangle)$. Once again, I will abuse notation for clarity.

3.1.1 Type System

I will now introduce the $\lambda_{\langle\text{op}\rangle}$ type system, by first introducing the types, and then the typing rules. The $\lambda_{\langle\text{op}\rangle}$ types are summarised in Figure 3.3. I highlight three important details: types are stratified into two levels (-1 for compile-time and 0 for run-time), effect rows are similarly stratified, and run-time code is made available at compile-time via a Code type.

First, **types are stratified into two levels, T^0 (run-time), and T^{-1} (compile-time).**

To motivate this stratification, consider the following question: what is the type of the number 3 in $\lambda_{\langle\text{op}\rangle}$? Perhaps surprisingly, the answer is not \mathbb{N} . Since we are working with

a two-stage system, we must be careful to disambiguate between run-time naturals and compile-time naturals, since these are not interchangeable. For example, the following program should **not** be well-typed, since 3 is a compile-time natural, whereas x is a run-time natural.

$$\lambda x : \mathbb{N}. \$ (3 + x)$$

However, removing the splice makes the program well-typed

$$\lambda x : \mathbb{N}. 3 + x$$

Following Xie et al. [35], I introduce integer levels to enforce separation between compile-time and run-time naturals. While the precise notion of level is slightly more involved (see below), for my purposes, it is sufficient to think of level 0 as run-time (so \mathbb{N}^0 is a run-time natural), and -1 for compile-time. The ill-typed example becomes

$$\lambda x : \mathbb{N}^0. \$ ((3 : \mathbb{N}^{-1}) + (x : \mathbb{N}^0))$$

and the well-typed example

$$\lambda x : \mathbb{N}^0. (3 : \mathbb{N}^0) + (x : \mathbb{N}^0)$$

More precisely, levels are defined as follows:

Definition 3.1.1 (Level) *The level of an expression e is calculated by subtracting the number of surrounding splices from the number of surrounding quotations.*

The definition of level generalises to multi-stage languages, where negative levels $(-1, -2, \dots)$ represent compile-time and non-negative levels $(0, 1, \dots)$ represent run-time. In a multi-staged language, separation is even more granular: for example, level 1 and level 0 naturals, despite both being run-time naturals, are disambiguated. However, since we only deal with two stages, we only consider two levels, 0 and -1 . The definition, and the examples above, further imply that the “default” level, in the absence of quotes and splices, is level 0. Intuitively, in the absence of quotes and splices, the programmer is ignoring metaprogramming facilities, and constructing a run-time program.

Notice that the opening question was slightly devious¹! We cannot assign a type to *program fragments*, like 3, since without knowledge of the wider context, we cannot know which level we are at: in the ill-typed example, 3 occurs under a splice, but no quotes, so it has type \mathbb{N}^{-1} , and in the well-typed example, it has type \mathbb{N}^0 . Unless otherwise stated, I will always assume program fragments are not nested in any quotes or splices, and thus occur at level 0.

Second, **effect rows are stratified into ξ (run-time) and Δ (compile-time)**.

In the following example, we print 1 at compile-time, and 2 at run-time. Further, we read an integer at run-time.

$$\$ (\text{print}(1); \langle\langle \text{print}(2); \text{readInt}() \rangle\rangle)$$

Hence, $\Delta = \{\text{print}\}$ and $\xi = \{\text{print}, \text{readInt}\}$. We may now disambiguate between different computation types:

¹sorry


3 Calculus

T^0	Compile-time value, run-time value (value types) <i>Example:</i> The type of x in $\lambda x.\text{return } x$
$T^0! \xi$	Compile-time value, run-time computation <i>Example:</i> The type of x in $\$(\text{do } x \leftarrow \langle\langle \text{return } 1 \rangle\rangle \text{ in return } x)$
$T^0! \Delta$	Compile-time computation, run-time value <i>Example:</i> $\lambda x.\text{return } x$
$T^0! \Delta; \xi$	Compile and run-time computation <i>Example:</i> $\$(\text{do } x \leftarrow \langle\langle \text{return } 1 \rangle\rangle \text{ in return } x)$

Third, there is an level -1 **Code type, representing run-time ASTs.**

By stratifying types to two levels, we have ensured that run-time (resp. compile-time) terms only interact with run-time (compile-time) terms. However, to enable meta-programming, run-time terms *should be available* at compile-time as ASTs. This is exactly the role of the Code type, thus allowing level -1 programs to manipulate ASTs of level 0 terms.

Putting it all together, we can now interpret complex $\lambda_{\langle\text{op}\rangle}$ types, like



Typing Judgement

Having described the types, I now present the type system. The typing rules are collated in [Figures 3.5](#) and [3.6](#). I will first explain the typing judgement, then highlight some key rules.

The shape of the typing judgement is mostly familiar, though, as in Xie et al. [35], I add level information and compiler modes. Level information will turn out to be redundant, but we will revisit this later.

$$\Gamma \vdash_{\text{Mode}}^{\text{Level}} e : T$$

Level. Recall that, when describing the λ_{op} types, I argued that one cannot type a program fragment, like 3, directly. One must also know the *level* (0 or -1), which is accordingly attached to the typing judgement.

Mode. For the purposes of elaboration, it can be useful to classify code into three categories:

- c** Code that is **ambient** and **inert**.
No surrounding quotes or splices
- s** Code that **manipulates ASTs** at compile-time.
Last surrounding annotation is a splice
- q** Code that **builds ASTs** to be manipulated at compile time.
Last surrounding annotation is a quote




Figure 3.4: Transitions between modes **c**, **s**, and **q**. Top-level splices transition from **c** to **s**, quotes transition from **s** to **q**, and splices (under quotes) transition from **q** to **s**.

To illustrate the purpose of the modes, consider the following meta-program, which evaluates to the AST of $\lambda x.1 + 2 + 3$. This program adopts the shorthand of [Section 2.2.2](#), where **returns** are implicitly inserted (for example, $\langle\langle 1 \rangle\rangle$ should really be $\langle\langle \text{return } 1 \rangle\rangle$), and we elide the order of operation ($1 + 2$ should really be written out using **do**). The emphasis is on the three modes:

$\lambda x. \$(\text{do } f \leftarrow (\lambda y. \langle\langle \$ (y) + 2 \rangle\rangle) \text{ in do } a \leftarrow \langle\langle 1 \rangle\rangle \text{ in } fa) + 3$
c s s q s q s c

- c** Identifies AST nodes that are **ambient** (within which computation may take place) and **inert** (cannot themselves be manipulated at compile-time).

$\lambda x. \$(\text{do } f \leftarrow (\lambda y. \langle\langle \$ (y) + 2 \rangle\rangle) \text{ in do } a \leftarrow \langle\langle 1 \rangle\rangle \text{ in } fa) + 3$
c s s q s q s c

- s** Identifies code that can be executed at compile-time to **manipulate ASTs**. Will be fully reduced at compile-time, and will not appear at run-time.

$\lambda x. \$(\text{do } f \leftarrow (\lambda y. \langle\langle \$ (y) + 2 \rangle\rangle) \text{ in do } a \leftarrow \langle\langle 1 \rangle\rangle \text{ in } fa) + 3$
c s s q s q s c

- q** Identifies code that **builds ASTs** (like **c**-mode) that can be manipulated at compile-time (unlike **c**-mode), to create run-time programs.

$\lambda x. \$(\text{do } f \leftarrow (\lambda y. \langle\langle \$ (y) + 2 \rangle\rangle) \text{ in do } a \leftarrow \langle\langle 1 \rangle\rangle \text{ in } fa) + 3$
c s s q s q s c

We may also describe how *transitions* between modes occur:

1. Top-level splices ($\$e$) transition from **c** (outside the splice) to **s** (within).
2. Quotes ($\langle\langle e \rangle\rangle$) transition from **s** (outside the quote) to **q** (within).
3. Splices ($\$e$) transition from **q** (outside the splice) to **s** (within).

Transitions between modes are illustrated in [Figure 3.4](#).

Since $\lambda_{\langle\text{op}\rangle}$ has only two levels, and my type system will ban nested splices and quotations ($\$ \e and $\langle\langle\langle\langle e \rangle\rangle\rangle\rangle$ are not valid program fragments), the compiler mode will uniquely identify the level (**c** and **q** imply level 0, and **s** level -1). I thus drop the level from my typing judgement, leaving it implicit.

Further, the typing judgements for **c** and **q** will be identical in almost all cases. To avoid repetition, I introduce the following notation

$$\Gamma \vdash_{\mathbf{c|q}} e : T$$

Where, for example, the typing rule

$$\frac{\Gamma_1 \vdash_{\mathbf{c|q}} e_1 : T_1 \quad \cdots \quad \Gamma_n \vdash_{\mathbf{c|q}} e_n : T_n}{\Gamma \vdash_{\mathbf{c|q}} e : T}$$

stands for two typing rules, one in **c**-mode and one in **q**-mode.

$$\begin{array}{c} \Gamma_1 \vdash_{\mathbf{c}} e_1 : T_1 \\ \vdots \\ \Gamma_n \vdash_{\mathbf{c}} e_n : T_n \\ \hline \Gamma \vdash_{\mathbf{c}} e : T \end{array} \qquad \begin{array}{c} \Gamma_1 \vdash_{\mathbf{q}} e_1 : T_1 \\ \vdots \\ \Gamma_n \vdash_{\mathbf{q}} e_n : T_n \\ \hline \Gamma \vdash_{\mathbf{q}} e : T \end{array}$$

The **c** and **q**-mode typing rules are summarised in Figure 3.5. The **s**-mode typing rules are summarised in Figure 3.6. In all modes, rules are extremely similar to the λ_{op} typing rules (Figure 2.4, Page 17). I focus on three important rules: **s**-QUOTE, **q**-SPLICE, and **c**-SPLICE.

Recall that the Code type makes available at compile-time a representation of ASTs of level 0 programs. Recall further that $\langle\langle e \rangle\rangle$ is the mechanism for *creating* ASTs (elements of type Code) from run-time programs. This intuition is captured by the **s**-QUOTE rule, where, to verify that $\langle\langle e \rangle\rangle$ is a valid AST of type Code, we verify that e is a run-time program of the corresponding type.

$$\begin{array}{c} (\mathbf{s}\text{-QUOTE}) \\ \Gamma \vdash_{\mathbf{q}} e : T^0 ! \Delta; \xi \\ \hline \Gamma \vdash_{\mathbf{s}} \langle\langle e \rangle\rangle : \text{Code}(T^0 ! \xi)^{-1} ! \Delta \end{array}$$

The dual to $\langle\langle e \rangle\rangle$ is $\$e$, which *eliminates* compile-time ASTs by transforming them (back) into run-time code. This intuition is captured by the **q**-SPLICE and **c**-SPLICE rules, where, to verify that $\$e$ is a valid run-time program, we check that e is a valid compile-time AST of the corresponding type.

$$\begin{array}{c} (\mathbf{c}\text{-SPLICE}) \\ \Gamma \vdash_{\mathbf{s}} e : \text{Code}(T^0 ! \xi)^{-1} ! \Delta \\ \hline \Gamma \vdash_{\mathbf{c}} \$e : T^0 ! \Delta; \xi \end{array} \qquad \begin{array}{c} (\mathbf{q}\text{-SPLICE}) \\ \Gamma \vdash_{\mathbf{s}} e : \text{Code}(T^0 ! \xi)^{-1} ! \Delta \\ \hline \Gamma \vdash_{\mathbf{q}} \$e : T^0 ! \Delta; \xi \end{array}$$

Note that since there are no **q**-QUOTE or **s**-SPLICE rules, the type system bans nested splices and quotations, and thus we can focus purely on levels 0 and -1 .

Finally, I define the notion of a well-typed $\lambda_{\langle\text{op}\rangle}$ expression.

Definition 3.1.2 (Well-typed expression) e is well-typed if $\vdash_{\mathbf{c}} e : T^0 ! \emptyset; \emptyset$

c and q Typing Rules $\lambda_{\langle\text{op}\rangle}$

(NAT)	(VAR)	(LAMBDA)
$\frac{}{\Gamma \vdash_{\text{c q}} m : \mathbb{N}^0! \Delta}$	$\frac{\Gamma(x) = T^0}{\Gamma \vdash_{\text{c q}} x : T^0! \Delta}$	$\frac{\Gamma, x : T_1^0 \vdash_{\text{c q}} e : T_2^0! \Delta; \xi}{\Gamma \vdash_{\text{c q}} \lambda x. e : (T_1^0 \xrightarrow{\xi} T_2^0)^0! \Delta}$
(APP)	(CONTINUE)	
$\frac{\Gamma \vdash_{\text{c q}} v_1 : (T_1^0 \xrightarrow{\xi} T_2^0)^0! \Delta \quad \Gamma \vdash_{\text{c q}} v_2 : T_1^0! \Delta}{\Gamma \vdash_{\text{c q}} v_1 v_2 : T_2^0! \Delta; \xi}$	$\frac{\Gamma \vdash_{\text{c q}} v_1 : (T_1^0 \xrightarrow{\xi} T_2^0)^0! \Delta \quad \Gamma \vdash_{\text{c q}} v_2 : T_1^0! \Delta}{\Gamma \vdash_{\text{c q}} \text{continue } v_1 v_2 : T_2^0! \Delta; \xi}$	
(RETURN)	(DO)	
$\frac{\Gamma \vdash_{\text{c q}} v : T^0! \Delta}{\Gamma \vdash_{\text{c q}} \text{return } v : T^0! \Delta; \xi}$	$\frac{\Gamma \vdash_{\text{c q}} e_1 : T_1^0! \Delta; \xi \quad \Gamma, x : T_1^0 \vdash_{\text{c q}} e_2 : T_2^0! \Delta; \xi}{\Gamma \vdash_{\text{c q}} \text{do } x \leftarrow e_1 \text{ in } e_2 : T_2^0! \Delta; \xi}$	
(OP)	(HANDLE)	
$\frac{\Gamma \vdash_{\text{c q}} v : T_1^0! \Delta \quad \text{op} : T_1^0 \rightarrow T_2^0 \in \Sigma \quad \text{op} \in \xi}{\Gamma \vdash_{\text{c q}} \text{op}(v) : T_2^0! \Delta; \xi}$	$\frac{\Gamma \vdash_{\text{c q}} e : T_1^0! \Delta; \xi \quad \Gamma \vdash_{\text{c q}} h : (T_1^0! \xi \Longrightarrow T_2^0! \xi')! \Delta \quad \forall \text{op} \in \Delta \setminus \Delta' . \text{op} \in \text{dom}(h)}{\Gamma \vdash_{\text{c q}} \text{handle } e \text{ with } \{h\} : T_2^0! \Delta; \xi'}$	
(RET-HANDLER)		
$\frac{\Gamma, x : T_1^0 \vdash_{\text{c q}} e : T_2^0! \Delta; \xi'}{\Gamma \vdash_{\text{c q}} \text{return}(x) \mapsto e : (T_1^0! \xi \Longrightarrow T_2^0! \xi')! \Delta}$		
(OP-HANDLER)		
$\frac{\text{op} : A^0 \rightarrow B^0 \in \Sigma \quad \Gamma \vdash_{\text{c q}} h : (T_1^0! \xi \Longrightarrow T_2^0! \xi')! \Delta \quad \Gamma, x : A^0, k : (B^0 \xrightarrow{\xi'} T_2^0)^0 \vdash_{\text{c q}} e : T_2^0! \Delta; \xi' \quad \xi \subseteq \xi \setminus \{\text{op}\} \quad \text{op}(x', k') \mapsto e' \notin h}{\Gamma \vdash_{\text{c q}} h; \text{op}(x, k) \mapsto e : (T_1^0! \xi \Longrightarrow T_2^0! \xi')! \Delta}$		
(c-SPLICE)	(q-SPLICE)	
$\frac{\Gamma \vdash_{\text{s}} e : \text{Code}(T^0! \xi)^{-1}! \Delta}{\Gamma \vdash_{\text{c}} \$e : T^0! \Delta; \xi}$	$\frac{\Gamma \vdash_{\text{s}} e : \text{Code}(T^0! \xi)^{-1}! \Delta}{\Gamma \vdash_{\text{q}} \$e : T^0! \Delta; \xi}$	

Figure 3.5: The **c**-mode and **q**-mode typing rules for $\lambda_{\langle\text{op}\rangle}$. The rules are nearly identical to the λ_{op} typing rules, with the exception of level annotations on types. Two additional rules, (**c**-SPLICE) (top-level splice) and (**q**-SPLICE) formalise the transition to **s**-mode.

s Typing Rules

$\text{(s-NAT)} \quad \frac{}{\Gamma \vdash_{\mathbf{s}} m : \mathbb{N}^{-1}}$	$\text{(s-VAR)} \quad \frac{\Gamma(x) = T^{-1}}{\Gamma \vdash_{\mathbf{s}} x : T^{-1}}$
$\text{(s-FUNCTION)} \quad \frac{\Gamma, x : T_1^{-1} \vdash_{\mathbf{s}} e : T_2^{-1} ! \Delta}{\Gamma \vdash_{\mathbf{s}} \lambda x.e : (T_1^{-1} \multimap T_2^{-1})^{-1}}$	$\text{(s-CONTINUATION)} \quad \frac{\Gamma, x : T_1^{-1} \vdash_{\mathbf{s}} e : T_2^{-1} ! \Delta}{\Gamma \vdash_{\mathbf{s}} \kappa x.e : (T_1^{-1} \multimap T_2^{-1})^{-1}}$
$\text{(s-APP)} \quad \frac{\Gamma \vdash_{\mathbf{s}} v_1 : (T_1^{-1} \multimap T_2^{-1})^{-1} \quad \Gamma \vdash_{\mathbf{s}} v_2 : T_1^{-1}}{\Gamma \vdash_{\mathbf{s}} v_1 v_2 : T_2^{-1} ! \Delta}$	$\text{(s-CONTINUE)} \quad \frac{\Gamma \vdash_{\mathbf{s}} v_1 : (T_1^{-1} \multimap T_2^{-1})^{-1} \quad \Gamma \vdash_{\mathbf{s}} v_2 : T_1^{-1}}{\Gamma \vdash_{\mathbf{s}} \text{continue } v_1 v_2 : T_2^{-1} ! \Delta}$
$\text{(s-RETURN)} \quad \frac{\Gamma \vdash_{\mathbf{s}} v : T^{-1}}{\Gamma \vdash_{\mathbf{s}} \text{return } v : T^{-1} ! \Delta}$	$\text{(s-DO)} \quad \frac{\Gamma \vdash_{\mathbf{s}} e_1 : T_1^{-1} ! \Delta \quad \Gamma, x : T_1 \vdash_{\mathbf{s}} e_2 : T_2^{-1} ! \Delta}{\Gamma \vdash_{\mathbf{s}} \text{do } x \leftarrow e_1 \text{ in } e_2 : T_2^{-1} ! \Delta}$
$\text{(s-OP)} \quad \frac{\Gamma \vdash_{\mathbf{s}} v : T_1^{-1} \quad \text{op} : T_1^{-1} \rightarrow T_2^{-1} \in \Sigma \quad \text{op} \in \Delta}{\Gamma \vdash_{\mathbf{s}} \text{op}(v) : T_2^{-1} ! \Delta}$	$\text{(s-HANDLE)} \quad \frac{\Gamma \vdash_{\mathbf{s}} e : T_1^{-1} ! \Delta \quad \Gamma \vdash_{\mathbf{s}} h : (T_1^{-1} ! \Delta \Longrightarrow T_2^{-1} ! \Delta')^{-1} \quad \forall \text{op} \in \Delta \setminus \Delta'. \text{op} \in \text{dom}(h)}{\Gamma \vdash_{\mathbf{s}} \text{handle } e \text{ with } \{h\} : T_2^{-1} ! \Delta'}$
$\text{(s-RET-HANDLER)} \quad \frac{\Gamma, x : T_1^{-1} \vdash_{\mathbf{s}} e : T_2^{-1} ! \Delta'}{\Gamma \vdash_{\mathbf{s}} \text{return}(x) \mapsto e : (T_1^{-1} ! \Delta \Longrightarrow T_2^{-1} ! \Delta')^{-1}}$	
$\text{(s-OP-HANDLER)} \quad \frac{\text{op} : A^{-1} \rightarrow B^{-1} \in \Sigma \quad \Gamma \vdash_{\mathbf{s}} h : (T_1^{-1} ! \Delta \Longrightarrow T_2^{-1} ! \Delta')^{-1} \quad \Gamma, x : A^{-1}, k : (B^{-1} \multimap T_2^{-1})^{-1} \vdash_{\mathbf{s}} e : T_2^{-1} ! \Delta' \quad \Delta' \subseteq \Delta \setminus \{\text{op}\} \quad \text{op}(x', k') \mapsto e' \notin h}{\Gamma \vdash_{\mathbf{s}} h; \text{op}(x, k) \mapsto e : (T_1^{-1} ! \Delta \Longrightarrow T_2^{-1} ! \Delta')^{-1}}$	
$\text{(s-QUOTE)} \quad \frac{\Gamma \vdash_{\mathbf{q}} e : T^0 ! \Delta; \xi}{\Gamma \vdash_{\mathbf{s}} \langle \langle e \rangle \rangle : \text{Code}(T^0 ! \Delta; \xi)^{-1} ! \Delta}$	

Figure 3.6: The **s**-mode typing rules for $\lambda_{\langle \text{op} \rangle}$. The rules are exactly identical to the λ_{op} typing rules, with the exception of level annotations on types, and the additional (QUOTE) rule, which makes level 0 code available at compile-time.


3.2 The Core Language: $\lambda_{\text{AST}(\text{op})}$

I now describe the core language, $\lambda_{\text{AST}(\text{op})}$, which is a simple extension of λ_{op} . For clarity, I rename λ_{op} values (v) and computations (c) to $\lambda_{\text{AST}(\text{op})}$ normal forms (n) and terms (t) respectively.


To construct $\lambda_{\text{AST}(\text{op})}$, I extend λ_{op} in two ways. First, I add machinery for metaprogramming:

1. **AST** nodes (like `Nat` and `Var`) and **binders** (`Var`)

The syntax explicitly disambiguates between binders (`Var`) and AST nodes. For example, consider the AST of $\lambda\alpha:\mathbb{N}. \alpha$, where I contrast the binder (left subtree) with the usage of the binder (right subtree).



This syntactic distinction is important because Binders and ASTs need to be distinguished at the type level, and I wish to keep the typing rules syntax directed. To see why it is important to distinguish Binders and ASTs at the *type* level, consider the following malformed AST:



The aforementioned malformed AST should be ill-typed. To do so, it is insufficient to require that the left sub-tree is of type $\text{AST}(\mathbb{N})$. Thus, Binders and ASTs should be distinguished at the type level.

2. **binderToAST**, a primitive for turning binders (`Var`) into *usages* of binders (`Var`)

This allows us to write programs like the following:

```
do x ← Var( $\alpha_{\mathbb{N}}$ ) in
do body ← binderToAST x in
return Lam(x, body)
```

 $\lambda_{\text{AST}(\text{op})}$

```
return Lam(Var( $\alpha_{\mathbb{N}}$ ), Var( $\alpha_{\mathbb{N}}$ ))
```

3. **mkvar** T , a primitive for generating fresh binders of type T (`Var`)

To illustrate why **mkvar** is necessary, recall that $\lambda_{\text{AST}(\text{op})}$ acts as an elaboration target for $\lambda_{\langle\text{op}\rangle}$. Consider the following $\lambda_{\langle\text{op}\rangle}$ program, which should ideally generate the AST of $\lambda\alpha.\lambda\beta. \text{return } (\alpha, \beta)$

```
$(do mkfun ← λk. ⟨⟨ λx : ℕ⁰. $(k ⟨x⟩) ⟩⟩ in
  mkfun ( λa. mkfun (λb. return (a, b)) ) )
```

 $\lambda_{\langle \text{op} \rangle}$

```
return Lam(Var(αℕ), Lam(Var(βℕ), Ret(Pair(Var(αℕ), Var(βℕ))))
```

The compile-time function `mkfun` is a higher-order function that takes some compile-time function k . k takes in a binder, x , and constructs the body of a function $\lambda x.\text{body}$. For example, in order to generate the body $x + 1$, we could write

```
mkfun (λa.⟨⟨$(a) + 1⟩⟩)
```

In the example $\lambda_{\langle \text{op} \rangle}$ program, k calls `mkfun`. This means that one constructs a *nested* function, whose formal parameters are bound to a and b respectively. If we simply elaborated x into $\text{Var}(x_{\mathbb{N}})$, we would bind both a and b to $\text{Var}(x_{\mathbb{N}})$, and generate the following (incorrect) AST:

```
return Lam(Var(xℕ), Lam(Var(xℕ), Ret(Pair(Var(xℕ), Var(xℕ))))
```

The problem is that we do not have α -renaming “for free”. We thus need to ensure that `mkfun` is elaborated into a function that generates *fresh* names for x each time it is called. This is the purpose of `mkvar`.

Second, I extend λ_{op} with machinery for scope extrusion checking. This comprises:

1. **err**, an error state for indicating the presence of scope extrusion,
2. **check**, a guarded **return** construct that either detects scope extrusion, and transitions to **err**, or does not detect scope extrusion, and transitions to **return**,
3. **check_M**, which behaves similarly to **check**, but (for reasons I will explain in [Section 4.3](#)) allows a set of *muted* variables M to *temporarily* extrude their scope,
4. **dlet**, a primitive for tracking which variables are well-scoped and which have extruded their scope,
5. **tls**, a marker for where top-level splices would have occurred in the source program. This is for *ease of reasoning only*.

Notice that, while the calculus the *machinery* for scope extrusion checking, it does not demand that one *use* it, or use it *properly*. Scope extrusion checking is not a language feature, but an algorithm one builds on top of the calculus.

3.2.1 Operational Semantics

I will now describe the operational semantics of $\lambda_{\text{AST}(\text{op})}$. Many rules are identical to those of λ_{op} ([Figure 2.2](#)): interesting rules are collated in [Figure 3.8](#). Rules are divided into those related to AST construction (**AST-RULE**), and those related to scope extrusion checking (**SEC-RULE**). I will now explain key rules.

Syntax

 $\lambda_{\text{AST}(\text{op})}$

Normal Forms	$n ::= \dots \mid \text{Nat}(m) \mid \text{Var}(\alpha_A) \mid \underline{\text{Var}}(\alpha_A) \mid \text{Lam}(n_1, n_2) \mid \text{App}(n_1, n_2) \mid \text{Continue}(n_1, n_2) \mid \text{Ret}(n) \mid \text{Do}(n_1, n_2, n_3) \mid \text{Op}(n) \mid \text{Hwith}(n_1, n_2) \mid \text{Hret}(n_1, n_2)(n_1, n_2) \mid \text{Hop}(n_1, n_2, n_3, n_4)$
Terms	$t ::= \dots \mid \text{check } n \mid \text{check}_M n \mid \text{dlet}(n, t) \mid \text{tls}(t) \mid \text{err} \mid \text{binderToAST } n$

Figure 3.7: $\lambda_{\text{AST}(\text{op})}$ syntax. The syntax is broadly the same as λ_{op} , except with the addition of AST constructors and scope extrusion checking machinery.

Configurations

Like λ_{op} , the operational semantics is defined over *configurations*. In $\lambda_{\text{AST}(\text{op})}$, configurations have the form:

$$\langle t; E; U; M; I \rangle$$

I will describe each element of the configuration at a high level: their roles will become clearer as we introduce each rule. t and E are as they are in λ_{op} : terms and evaluation contexts. U acts as a source of freshness for name generation. M is a set of muted variables, i.e. those that we do not want to trigger a scope extrusion error, even if they have extruded their scope. I indicates the point at which we should *unmute*, setting M to \emptyset .

AST Rules

The **AST-GEN** rule describes the behaviour of **mkvar**: **mkvar** T produces a **Var** of type T and some name α . Recall that names should be **fresh**: that is, multiple calls to **mkvar** should always return variables with different names. In order to ensure that names are fresh, we need to keep track of names that have been previously generated. This is the purpose of U in the configuration, and the side condition on the rule. To ensure determinacy of the semantics, we will assume that names are chosen *deterministically*.

The other primitive, **binderToAST**, turns binders $\underline{\text{Var}}(\alpha_T)$ into usages of the binder $\text{Var}(\alpha_T)$ (**AST-USE**).

Scope Extrusion Checking Rules

More interesting are the primitives for scope extrusion checking. The **check** primitive acts like a guarded **return**, which can catch occurrences of scope extrusion. For some arbitrary n of AST type, either:

1. All the free variables of n are properly scoped, so **check** n reduces to **return** n (**SEC-CHS**)
2. Some free variables of n are not properly scoped, so **check** n reduces to **err** (**SEC-CHF**)

Operational Semantics

Selected Rules

Evaluation Contexts

$$F ::= \dots \mid \mathbf{dlet}(\underline{\text{Var}}(\alpha_A), [-]) \mid \mathbf{tls}([-])$$

Operational Semantics

(AST-SYM)	$\mathbf{mkvar} \ A; E; U; M; I$	\rightarrow	$\mathbf{return} \ \underline{\text{Var}}(\alpha_A); E; U \cup \{\alpha\}; M; I$ (where $\alpha \notin U$)
(AST-USE)	$\mathbf{binderToAST} \ \underline{\text{Var}}(\alpha_A); E; U; M; I$	\rightarrow	$\mathbf{return} \ \text{Var}(\alpha_A); E; U; M; I$
(SEC-CHS)	$\mathbf{check} \ n; E; U; M; I$	\rightarrow	$\mathbf{return} \ n; E; U; M; I$ (if $\text{FV}^0(n) \subseteq \pi_{\text{Var}}(E)$)
(SEC-CHF)	$\mathbf{check} \ n; E; U; M; I$	\rightarrow	$\mathbf{err}; E; U; M; I$ (if $\text{FV}^0(n) \not\subseteq \pi_{\text{Var}}(E)$)
(SEC-CMS)	$\mathbf{check}_M \ n; E; U; M; I$	\rightarrow	$\mathbf{return} \ n; E; U; M; I$ (if $\text{FV}^0(n) \setminus M \subseteq \pi_{\text{Var}}(E)$)
(SEC-CMF)	$\mathbf{check}_M \ n; E; U; M; I$	\rightarrow	$\mathbf{err}; E; U; M; I$ (if $\text{FV}^0(n) \setminus M \not\subseteq \pi_{\text{Var}}(E)$)
(SEC-TLS)	$\mathbf{tls}(\mathbf{return} \ n); E; U; M; I$	\rightarrow	$\mathbf{return} \ n; E; U; M'; I'$
(SEC-DLT)	$\mathbf{dlet}(\underline{\text{Var}}(\alpha_T), \mathbf{return} \ n); E; U; M; I$	\rightarrow	$\mathbf{return} \ n; E; U; M'; I'$ (if $I < \text{len}(E)$ then $M' = M, I' = I$ else $M' = \emptyset, I' = \top$)
(EFF-OP)	$\mathbf{op}(v); E_1[\mathbf{handle} \ E_2 \ \mathbf{with} \ \{h\}]; U; M; I$	\rightarrow	$c[v/x, \text{cont}/k]; E_1; U; M \cup \pi_{\text{Var}}(E_2); I'$ (where $\text{cont} = \kappa x. \mathbf{handle} \ E_2[\mathbf{return} \ x] \ \mathbf{with} \ \{h\}$ and $\mathbf{op}(x, k) \mapsto c \in h$ and $\mathbf{op} \notin \text{handled}(E_2)$ and $I' = \min(\text{len}(E), I)$)

Figure 3.8: Selected rules of the $\lambda_{\text{AST}(\text{op})}$ operational semantics. Many of the rules can be trivially adapted from the λ_{op} semantics (Figure 2.2). The muting and unmuting of variables is complex, and will be best explained when we discuss scope extrusion checks. For now, these mechanisms are **highlighted**.

What does it mean to be “properly scoped” (the side condition on SEC-CHS)? The answer is slightly subtle. Consider the following program

do body \leftarrow **check** n **in** **check** $\text{Lam}(\text{Var}(\alpha_{\mathbb{N}}), \text{body})$

I argue that $\text{Var}(\alpha_{\mathbb{N}})$ **ought to be** properly scoped in n (should not cause a transition to **err**). However, it is hard to deduce this from the *static* structure of the program. Instead, one has to reason about the *dynamic* execution of the program. Rather than calculating what is properly scoped as a *language feature*, I defer it to the programmer. The programmer must *declare* that a variable is properly scoped through use of the **dlet** keyword.

dlet($\underline{\text{Var}}(\alpha_{\mathbb{N}})$, **do** body \leftarrow **check** n **in** **check** $\text{Lam}(\text{Var}(\alpha_{\mathbb{N}}), \text{body})$)

More precisely, **dlet** places a frame of the form $\mathbf{dlet}(\underline{\text{Var}}(\alpha_A), [-])$ on the evaluation context E . I use the notation $\pi_{\text{Var}}(E)$ to filter out variables declared in this manner. For example,

$$\pi_{\text{Var}}(\mathbf{dlet}(\underline{\text{Var}}(\alpha_A), \mathbf{do} \ x \leftarrow [-] \ \mathbf{in} \ t)) = \{\text{Var}(\alpha_A)\}$$

Given a term of the form **check** n in some evaluation context E , where n is an AST, **check** thus checks that the free Vars of n , written $\text{FV}^0(n)$, have all been properly declared, or to be precise, a subset of $\pi_{\text{Var}}(E)$.

It may seem lazy of me to define the semantics of **check** in such a way that places the burden onto the user. Recall, however, that $\lambda_{\text{AST}(\text{op})}$ is *not* meant to be programmed in directly. Rather, it acts as an elaboration target for $\lambda_{\langle\langle\text{op}\rangle\rangle}$, and I define the elaboration. Therefore, I am the (only) $\lambda_{\text{AST}(\text{op})}$ user, and the onus is on me to justify that my elaboration uses **check** appropriately.

I also introduce **check** _{M} as a variant of **check**. As I will explain in [Section 4.3](#), to design a good scope extrusion check, it is necessary to *mute* some variables, pretending that they are properly scoped. **check** _{M} will behave exactly like **check**, except that it will also pretend the muted variables M are properly scoped.

Similarly, when justifying the correctness of scope extrusion checks, it will be useful to remember the position of top-level splices in the $\lambda_{\langle\langle\text{op}\rangle\rangle}$ source program. This is the purpose of **tls**, which should be interpreted as a no-op (SEC-TLS).

The final two rules, SEC-DLT and EFF-OP, behave as normal, but additionally mute or unmute Vars. The operations of muting and unmuting are best explained in [THE NEXT CHAPTER]. For now, they are **highlighted**, and can be mostly ignored. At a high level, when an operation is performed, we mute some set of variables, and potentially update the point at which they should be unmuted. When we remove a declared variable, we additionally check if we ought to unmute variables (and do so if we should).

3.2.2 Type System

Extending the types is similarly straightforward. I add only two types: a Binder type and an AST type. The types are summarised in [Figure 3.9](#).

Typing Rules

I now describe a selection of $\lambda_{\text{AST}(\text{op})}$ typing rules ([Figure 3.10](#)). The rules are extremely straightforward: $\text{Var}(\alpha_A)$ is a Binder of type A , and $\text{Var}(\alpha_A)$ is an AST of type A . **mkvar** A is a computation that produces Binders of type A , and $\text{Lam}(n_1, n_2)$ is well-typed if n_1 is a Binder and n_2 an AST.

Notice that this type system does not guarantee that the resulting AST is *well-scoped*, for example, the following is well-typed:

$$\cdot \vdash \text{Var}(\alpha_A) : \text{AST}(A)$$

The typing rules for scope extrusion checks are even more straightforward: they are effectively invisible to the type system. The only complex case is **err**, which can be assigned any type in any context. **err** thus behaves similarly to **absurd** in Haskell, or in the λ -calculus extended with the empty type [26].

Finally, I define the notion of well-typed term.

Definition 3.2.1 (Well-typed term) A term t is well-typed if $\cdot \vdash t : T! \emptyset$

Types*Computation and Handler Types omitted* $\lambda_{\text{AST}(\text{op})}$ **Run-time Pre-types**Effects Row $\xi ::= \emptyset \mid \xi \cup \{\text{op}_i\}$

Value type $A ::= \mathbb{N}$

$ A_1 \xrightarrow{\xi} A_2$	functions
$ A_1 \xrightarrow{\xi} A_2$	continuations

Computation type $A! \xi$ Handler type $A_1! \xi \implies A_2! \xi'$ **Types**

Value type $T ::= \dots$

$ \text{Binder}(A)$	binders
$ \text{AST}(A)$	AST (value)
$ \text{AST}(A! \xi)$	AST (computation)
$ \text{AST}(A_1! \xi \implies A_2! \xi')$	AST (handler)

Figure 3.9: The types of $\lambda_{\text{AST}(\text{op})}$. $\lambda_{\text{AST}(\text{op})}$ types extend λ_{op} types with an AST type (for ASTs), and a Binder type

3.2.3 Implementation

The core calculus $\lambda_{\text{AST}(\text{op})}$ describes abstractly a concrete OCaml implementation. In the concrete OCaml implementation, we do not need to introduce **check**, **dlet**, and **err** as primitives. Rather, they can be encoded as a *mode of use* of effects and handlers.

1. **check** n is implemented by performing a FreeVar effect, that are passed the free variables of n
check_M n is similar, except there are also Mute and Unmute effects
2. **dlet**($\text{Var}(\alpha_T), t$) is implemented as a *handler* of the FreeVar effect, which subtracts $\text{Var}(\alpha_T)$ from the set of free variables, and either:
 - a) Resumes the continuation, if the set of free variables is now empty (all free variables properly declared/scoped)
 - b) Performs another FreeVar effect, to check that the remaining free variables are properly declared. Following a successful such check, the continuation may be resumed.
3. **err** is an unhandled FreeVar effect

3.3 Elaboration from $\lambda_{\langle\langle\text{op}\rangle\rangle}$ to $\lambda_{\text{AST}(\text{op})}$

Having described both $\lambda_{\langle\langle\text{op}\rangle\rangle}$ and $\lambda_{\text{AST}(\text{op})}$, I will now describe a simple elaboration from $\lambda_{\langle\langle\text{op}\rangle\rangle}$ to $\lambda_{\text{AST}(\text{op})}$. This elaboration will be *simple* – it will not insert any dynamic scope extrusion checks.

Typing Rules

Selected Rules

$\lambda_{\text{AST}(\text{op})}$

(BINDER)	(VARIABLE)	(MKVAR)
$\frac{}{\Gamma \vdash \underline{\text{var}}(\alpha_A) : \text{Binder}(A)}$	$\frac{}{\Gamma \vdash \text{var}(\alpha_A) : \text{AST}(A)}$	$\frac{}{\Gamma \vdash \mathbf{mkvar} A : \text{Binder}(A) ! \Delta}$
(BINDERToAST)	(LAMBDA-AST)	
$\frac{\Gamma \vdash n : \text{Binder}(A)}{\Gamma \vdash \mathbf{binderToAST} n : \text{AST}(A) ! \Delta}$	$\frac{\Gamma \vdash n_1 : \text{Binder}(A_1) \quad \Gamma \vdash n_2 : \text{AST}(A_2 ! \xi)}{\Gamma \vdash \mathbf{Lam}(n_1, n_2) : \text{AST}(A_1 \xrightarrow{\xi} A_2)}$	
(ERR)	(TLS)	(DLET)
$\frac{}{\Gamma \vdash \mathbf{err} : T ! \Delta}$	$\frac{\Gamma \vdash t : T ! \Delta}{\Gamma \vdash \mathbf{tls}(t) : T ! \Delta}$	$\frac{\Gamma \vdash n : \text{Binder}(A) \quad \Gamma \vdash t : T ! \Delta}{\Gamma \vdash \mathbf{dlet}(n, t) : T ! \Delta}$
(CHECK)		
$\frac{\Gamma \vdash n : T \quad T = \text{AST}(A) \vee \text{AST}(A ! \xi) \vee \text{AST}(A_1 ! \xi \implies A_2 ! \xi')}{\Gamma \vdash \mathbf{check} n : T ! \Delta}$		

Figure 3.10: Selected $\lambda_{\text{AST}(\text{op})}$ typing rules

The elaboration is defined on typing judgements: I elaborate a $\lambda_{\langle\text{op}\rangle}$ judgement to a $\lambda_{\text{AST}(\text{op})}$ judgement. To do so, I define four elaborations: on effect rows, types, contexts, and terms. As it will be clear from context which elaboration is being referred to, I will abuse notation and write $\llbracket - \rrbracket$ for all four.

Elaborating effect rows will just be the identity, that is

$$\begin{aligned} \llbracket \Delta \rrbracket &= \Delta \\ \llbracket \xi \rrbracket &= \xi \end{aligned}$$

and I will not touch on them any further.

3.3.1 Elaborating Types

To define elaboration of types, it will be convenient to refer to a helper function, *erase*, that *erases* all of the level annotations (and elaborates effect rows). It is easy to define inductively, and I do not give a formal definition, but rather an example.

$$\text{erase}((T_1^0 \xrightarrow{\xi} T_2^0)^0) = T_1 \xrightarrow{\llbracket \xi \rrbracket} T_2$$

I can now define elaboration of types easily. In a nutshell, level 0 types elaborate into AST types. Level -1 types elaborate into themselves (sans level annotations), except

for Code types, which elaborate into AST types.

$$\begin{aligned}
\llbracket T^0 \rrbracket &= \text{AST}(\text{erase}(T^0)) \\
\llbracket T^0 ! \xi \rrbracket &= \text{AST}(\text{erase}(T^0 ! \xi)) \\
\llbracket T^0 ! \Delta \rrbracket &= \text{AST}(\text{erase}(T^0)) ! \llbracket \Delta \rrbracket \\
\llbracket T^0 ! \Delta ; \xi \rrbracket &= \text{AST}(\text{erase}(T^0 ! \xi)) ! \llbracket \Delta \rrbracket \\
\llbracket (T_1^0 ! \xi \Longrightarrow T_2^0 ! \xi')^0 \rrbracket &= \text{AST}(\text{erase}((T_1^0 ! \xi \Longrightarrow T_2^0 ! \xi')^0)) \\
\llbracket T^{-1} \rrbracket &= \text{if } T \neq \text{Code}(T^0 ! \xi) \text{ then } \text{erase}(T^{-1}) \\
&\quad \text{else } \text{AST}(\text{erase}(T^0 ! \xi)) \\
\llbracket T^{-1} ! \Delta \rrbracket &= \llbracket T^{-1} \rrbracket ! \llbracket \Delta \rrbracket \\
\llbracket (T_1^{-1} ! \Delta \Longrightarrow T_2^{-1} ! \xi')^{-1} \rrbracket &= \text{AST}(\text{erase}((T_1^{-1} ! \Delta \Longrightarrow T_2^{-1} ! \xi')^{-1}))
\end{aligned}$$

3.3.2 Elaborating Contexts

Elaborating contexts is *slightly* subtle. Rather than mapping the elaboration function over all types in the context, we treat the level 0 types slightly differently, elaborating in Binder, rather than AST types.

$$\begin{aligned}
\llbracket \cdot \rrbracket &= \cdot \\
\llbracket \Gamma, x : T^0 \rrbracket &= \llbracket \Gamma \rrbracket, x : \text{Binder}(\text{erase}(T^0)) \\
\llbracket \Gamma, x : T^{-1} \rrbracket &= \llbracket \Gamma \rrbracket, x : \llbracket T^{-1} \rrbracket
\end{aligned}$$

To see why this is the case, notice that the only cases where the context Γ is extended with a level 0 variable $x : T^0$ occur in **c** or **q**. In these modes, we are building ASTs, and thus x must be an AST Binder.

3.3.3 Elaborating Terms

Elaborating terms is slightly more involved. To start, we will assume that we have annotated all binders with their types, for example

$$\lambda x : \mathbb{N}^0. e$$

The elaboration for terms is moderated by the **mode**: **c**, **q**, or **s**. Selected rules are collated in [Figure 3.11](#).

At a high level, in **c** and **q**-mode, one builds ASTs, calling **mkvar** when binders are encountered (see earlier discussion on the necessity of **mkvar**). Elaboration in **c** and **q**-modes do not differ particularly significantly, with the exception of the rule for splice, where in **c**, I insert the marker for the top-level splice, and in **q**, I do not. Elaboration in **c** and **q**-modes will be further distinguished when I extend the elaboration to insert scope extrusion checks. Elaboration in **s**-mode is effectively the identity.

3.3.4 Elaborating Typing Judgements

We may now elaborate full typing judgements, by elaborating each component in the judgement. For example, take the typing judgement for lambdas in **c**-mode.

$$\frac{\Gamma, x : T_1^0 \vdash_{\mathbf{c}} e : T_2^0 ! \Delta ; \xi}{\Gamma \vdash_{\mathbf{c}} \lambda x. e : (T_1^0 \xrightarrow{\xi} T_2^0)^0 ! \Delta}$$

Term Elaboration

Selected Rules

 $\lambda_{\langle \text{op} \rangle}$

$$\begin{aligned}
 \llbracket x \rrbracket_{\text{c}|\text{q}} &= \text{binderToAST } x \\
 \llbracket \lambda x : T^0. e \rrbracket_{\text{c}|\text{q}} &= \text{do } x \leftarrow \text{mkvar } \text{erase}(T^0) \text{ in do body} \leftarrow \llbracket e \rrbracket_{\text{c}|\text{q}} \text{ in return Lam}(x, \text{body}) \\
 \llbracket \$e \rrbracket_{\text{c}} &= \text{t1s}(\llbracket e \rrbracket_{\text{s}}) \\
 \llbracket \$e \rrbracket_{\text{q}} &= \llbracket e \rrbracket_{\text{s}} \\
 \llbracket x \rrbracket_{\text{s}} &= x \\
 \llbracket \lambda x : T^0. e \rrbracket_{\text{s}} &= \lambda x. \llbracket e \rrbracket_{\text{s}} \\
 \llbracket \langle e \rangle \rrbracket_{\text{s}} &= \llbracket e \rrbracket_{\text{q}}
 \end{aligned}$$

Figure 3.11: Selected term elaboration rules from $\lambda_{\langle \text{op} \rangle}$ to $\lambda_{\text{AST}(\text{op})}$. Elaboration is moderated by the compiler mode. In **c** and **q**, elaboration builds ASTs. In **s**-mode, elaboration is effectively the identity.

which we elaborate by applying the elaboration component-wise

$$\frac{\llbracket \Gamma, x : T_1^0 \rrbracket \vdash \llbracket e \rrbracket_{\text{c}} : \llbracket T_2^0 ! \Delta; \xi \rrbracket}{\llbracket \Gamma \rrbracket \vdash \llbracket \lambda x. e \rrbracket_{\text{c}} : \llbracket (T_1^0 \xrightarrow{\xi} T_2^0)^0 ! \Delta \rrbracket}$$

Letting $A_i = \text{erase}(T_i^0)$, and $\llbracket e \rrbracket_{\text{c}} = t$, and applying the elaboration functions defined above, we obtain

$$\frac{\llbracket \Gamma \rrbracket, x : \text{Binder}(A_1) \vdash t : \text{AST}(A_2 ! \xi) ! \Delta}{\llbracket \Gamma \rrbracket \vdash \text{do } x \leftarrow \text{mkvar } \text{erase}(T^0) \text{ in do body} \leftarrow t \text{ in return Lam}(x, \text{body}) : \text{AST}(A_1 \xrightarrow{\xi} A_2) ! \Delta}$$

which, assuming that the premise is valid typing derivation, corresponds to a valid $\lambda_{\langle \text{op} \rangle}$ typing derivation.

Is this true in general? Do $\lambda_{\langle \text{op} \rangle}$ typing derivations always elaborate into $\lambda_{\text{AST}(\text{op})}$ typing derivations? Yes, but the question begets a larger point: what properties can we claim about the calculus as defined? What metatheoretic results may we establish?

3.4 Metatheory

I now establish several metatheoretic properties of the $\lambda_{\langle \text{op} \rangle}$ calculus.

The first states that well-typed $\lambda_{\langle \text{op} \rangle}$ programs elaborate into well-typed $\lambda_{\text{AST}(\text{op})}$ programs.

Theorem 3.4.1 (Elaboration Preservation) *If $\Gamma \vdash_{\star} e : \tau$ then $\llbracket \Gamma \rrbracket \vdash \llbracket e \rrbracket_{\star} : \llbracket \tau \rrbracket$, where $\star = \text{c} \mid \text{q} \mid \text{s}$ and τ is a level 0 or level -1 value, computation, or handler type.*

The proof is by induction on the typing rules, as in the example in the previous section.

We can also prove that the core language $\lambda_{\text{AST}(\text{op})}$ satisfies appropriate progress and preservation properties.

Theorem 3.4.2 (Progress) *If $\Gamma \vdash E[t] : T ! \Delta$ then for all U, M, I either*

1. *t of the form **return** n and $E = [-]$,*
2. *t of the form **op**(v) for some $op \in \Delta$,*
3. *t of the form **err***
4. *$\exists t', E', U', M', I'$ such that $\langle t; E; U; M; I \rangle \rightarrow \langle t'; E'; U'; M'; I' \rangle$*

Note that I additionally have to consider the error state **err**. Progress is proved by induction over the typing derivation. Since I am extending λ_{op} , I only need to consider the augmented typing rules, all of which are extremely straightforward.

Theorem 3.4.3 (Reduction Preservation) *If*

1. $\Gamma \vdash E[t] : T$ *and*
2. $\langle t; E; U; M; I \rangle \rightarrow \langle t'; E'; U'; M'; I' \rangle$

then $\Gamma \vdash E'[t'] : T$

Proof proceeds by induction over the operational semantics. Once again, I only need to consider the augmented rules, which are very simple.

As a corollary, we get a notion of type safety.

Corollary 3.4.1 (Type Safety) *If $\cdot \vdash_{\text{c}} e : T^0 ! \emptyset; \emptyset$ then either*

1. $\langle \llbracket e \rrbracket_{\text{c}}; [-]; \emptyset; \emptyset; \top \rangle \rightarrow^{\omega}$,
2. $\langle \llbracket e \rrbracket_{\text{c}}; [-]; \emptyset; \emptyset; \top \rangle \rightarrow^* \langle \mathbf{err}; E; U; M; I \rangle$ *for some E, U, M, I , or*
3. $\langle \llbracket e \rrbracket_{\text{c}}; [-]; \emptyset; \emptyset; \top \rangle \rightarrow^* \langle \mathbf{return } n; [-]; U; M; I \rangle$ *for some U, M, I*

Importantly, this notion of type safety is weak. A system which always reports a scope extrusion error (**err**) would be type safe. Further, a system which never reports scope extrusion would be type safe as well. Due to the potential presence of scope extrusion, in the third case of [Corollary 3.4.1](#), we cannot additionally claim that the normal form n represents a well-typed λ_{op} program.

Additionally, underneath a top-level splice, splice and quotation are duals. That is,

Theorem 3.4.4 (Quote-Splice Duality)

$$\begin{aligned} \$\langle\langle e \rangle\rangle &=_{\text{q}} e \\ \langle\langle \$e \rangle\rangle &=_{\text{s}} e \end{aligned}$$

Where, in this context, $=_{\star}$ means “elaborates to contextually equivalent $\lambda_{\text{AST}(\text{op})}$ programs in \star mode”. Parameterising by the mode is necessary, since the mode will affect the result of elaboration. Indeed, we may prove something stronger: they elaborate to

the same $\lambda_{\text{AST}(\text{op})}$ program (contextual equivalence follows from reflexivity). Proof is by inspection of the definition of elaboration, where

$$\begin{aligned} \llbracket \$\langle e \rangle \rrbracket_{\mathbf{q}} = t &\iff \llbracket e \rrbracket_{\mathbf{q}} = t \\ \llbracket \langle \$e \rangle \rrbracket_{\mathbf{s}} = t &\iff \llbracket e \rrbracket_{\mathbf{s}} = t \end{aligned}$$

4 Scope Extrusion

In the previous chapter, I described $\lambda_{\langle\text{op}\rangle}$, a calculus for studying the interaction of effects and metaprogramming. In this chapter, I illustrate how $\lambda_{\langle\text{op}\rangle}$ can be used to precisely define scope extrusion, and evaluate scope extrusion solutions. I describe and evaluate four solutions:

1. A Lazy Dynamic Check, first described by Kiselyov [14] (Section 4.1).
2. An Eager Dynamic Check, which to the best of my knowledge is a faithful description of the current MetaOCaml check (Section 4.2).
3. A novel Best-Effort Dynamic Check, inspired by the evaluations of the lazy and eager checks, and which I argue occupies a goldilocks zone between expressiveness and efficiency (Section 4.3).
4. Refined Environment Classifiers, a static approach first described by Kiselyov et al. [17] (Section 4.4).

The three dynamic checks have additionally been implemented in MacoCaml.

Finally, drawing on my use of $\lambda_{\langle\text{op}\rangle}$ for describing and evaluating the four solutions, I briefly evaluate $\lambda_{\langle\text{op}\rangle}$'s ability to facilitate comparative evaluation of scope extrusion solutions (Section 4.5).

4.1 Lazy Dynamic Check

Recall that one definition of scope extrusion relates to the *result* of compile-time execution (Section 2.3.1, Page 22) [14].

I call this lazy scope extrusion, and define it as a property on $\lambda_{\text{AST}(\text{op})}$ configurations as follows

Definition 4.1.1 (Lazy Scope Extrusion) *A $\lambda_{\text{AST}(\text{op})}$ configuration of the form*

$$\langle t; E; U; M; I \rangle$$

exhibits lazy scope extrusion if all of the following hold:

1. $t = \text{return } n$ for some n of AST type
2. $E = E' :: \mathbf{tls}([-])$ for some E'
3. $FV^0(n) \not\subseteq \pi_{\text{var}}(E)$

Note that the marker for top-level splice \mathbf{tls} implies that we splice n into a larger, ambient and inert, program.

I now define the lazy dynamic check precisely as a modified term elaboration $\llbracket \cdot \rrbracket^{\text{Lazy}}$. I need to make two changes: first, inserting **dlets** into **c** cases, for example

$$\llbracket \lambda x : T^0. e \rrbracket_c^{\text{Lazy}} = \text{do } x \leftarrow \text{mkvar } \text{erase}(T^0) \text{ in } \text{dlet}(x, \text{do body} \leftarrow \llbracket e \rrbracket_c^{\text{Lazy}} \text{ in } \text{return } \text{Lam}(x, \text{body}))$$

second, to perform a check after a top-level splice:

$$\llbracket \$e \rrbracket_c^{\text{Lazy}} \triangleq \text{do } \text{res} \leftarrow \text{tls}(\llbracket e \rrbracket_s^{\text{Lazy}}) \text{ in } \text{check res}$$

Due to the simplicity of the algorithm, verifying the correctness and expressiveness of the check is trivial: the check reports scope extrusion if, and only if, the program exhibits lazy scope extrusion.

Theorem 4.1.1 (Correctness and Expressiveness of the Lazy Dynamic Check) *If*

1. $\cdot \vdash_c e : T^0 ! \emptyset; \emptyset$,
2. $\llbracket e \rrbracket_c^{\text{Lazy}} = t$

Then $\langle t; [-]; \emptyset; \emptyset; \top \rangle \rightarrow^ \langle \text{err}; E; U; M; I \rangle$ if, and only if, for some E'*
 $\langle t; [-]; \emptyset; \emptyset; \top \rangle \rightarrow^* \langle \text{return } n; E'; U; M; I \rangle$, *and*
 $\langle \text{return } n; E'; U; M; I \rangle$ *exhibits lazy scope extrusion.*

However, due again to its simplicity, the lazy dynamic check is inefficient and uninformative, and therefore not suitable for practical use [14] (Section 2.3.1, Page 22).

4.2 Eager Dynamic Check

Yet another definition of scope extrusion relates to *intermediate results* during compile-time execution (Section 2.3.1, Page 22). For example, Kiselyov [14] defines scope extrusion as (emphasis mine)

At any point during the evaluation, an occurrence of an open-code value with a free variable whose name is not dynamically bound is called scope extrusion.

I call this eager scope extrusion, and define it as a property on $\lambda_{\text{AST}(\text{op})}$ configurations as follows.

Definition 4.2.1 (Eager Scope Extrusion) *A $\lambda_{\text{AST}(\text{op})}$ configuration of the form*

$$\langle t; E; U; M; I \rangle$$

exhibits eager scope extrusion if all of the following hold:

1. $t = \text{return } n$ *for some n of AST type*
2. $FV^0(n) \not\subseteq \pi_{\text{var}}(E)$

Notice that lazy scope extrusion is a special case of eager scope extrusion.

It is possible to extend the lazy dynamic check into an eager dynamic check. In addition to the top-level splice check, and the **c**-mode **dlets**, I add **dlets** and **checks** in **q**-mode, for example

$$\llbracket \lambda x : T^0. e \rrbracket_{\mathbf{q}}^{\text{Eager}} = \text{do } x \leftarrow \text{mkvar } \text{erase}(T^0) \text{ in } \text{check}(\text{dlet}(x, \text{do } \text{body} \leftarrow \llbracket e \rrbracket_{\mathbf{q}}^{\text{Eager}} \text{ in } \text{return } \text{Lam}(x, \text{body})))$$

where **check** e is syntactic sugar for **do** $x \leftarrow e$ **in** **check** x .

As another example,

$$\llbracket v_1 v_2 \rrbracket_{\mathbf{q}}^{\text{Eager}} = \text{do } f \leftarrow \llbracket v_1 \rrbracket_{\mathbf{q}}^{\text{Eager}} \text{ in } \text{do } a \leftarrow \llbracket v_2 \rrbracket_{\mathbf{q}}^{\text{Eager}} \text{ in } \text{check App}(f, a)$$

notice how **return** $\text{App}(f, a)$ is replaced by **check** $\text{App}(f, a)$.

Intuitively, we insert checks whenever we build ASTs. Hence, assume that we reduce to a configuration that exhibits eager scope extrusion, and let the offending AST be n . The error will be detected and reported the next time n is used to build a bigger AST in an unsafe way.

As an example, consider the following $\lambda_{\langle \text{op} \rangle}$ program in Listing 13. Eager scope extrusion is caused, with the offending program fragment being **return** y , where y refers to the unbound variable $\text{Var}(x_{\mathbb{N}})$ ¹. Eager scope extrusion is not caught immediately. Rather, y is bound to z , and scope extrusion is caught when z is used to build a larger AST, $\langle \$z + 1 \rangle$.

```
 $\$(\text{do } z \leftarrow (\text{handle}(\langle \lambda x. \$(\text{extrude}(\langle \langle x \rangle \rangle); \langle \langle \text{return } x \rangle \rangle)))$ 
   $\text{with } \{\text{return}(u) \mapsto \langle \langle 0 \rangle \rangle; \text{extrude}(y, k) \mapsto \text{return } y\}$ 
   $\text{in } \langle \langle \$z + 1 \rangle \rangle)$ 
```

$\lambda_{\langle \text{op} \rangle}$

Listing 13: Illustrating the eager dynamic check: eager scope extrusion is caused by the **return** y expression. Scope extrusion is not immediately detected. Rather, y is bound to z . Scope extrusion is reported when z is used to build a larger AST: $\langle \langle \$z + 1 \rangle \rangle$

The eager dynamic check also checks at top-level splices, like the lazy dynamic check. Conceptually, a top-level splice builds a larger, ambient and inert, AST (Listing 14):

```
 $\text{do } z \leftarrow \$(\text{handle}(\langle \lambda x. \$(\text{extrude}(\langle \langle x \rangle \rangle); \langle \langle \text{return } x \rangle \rangle)))$ 
   $\text{with } \{\text{return}(u) \mapsto \langle \langle 0 \rangle \rangle; \text{extrude}(y, k) \mapsto \text{return } y\}$ 
   $\text{in } z + 1$ 
```

$\lambda_{\langle \text{op} \rangle}$

Listing 14: The eager dynamic check additionally checks at top-level splices, similarly to the lazy dynamic check.

To the best of my knowledge, the eager dynamic check is a faithful model of the MetaOCaml check, as described by Kiselyov [16].

4.2.1 Correctness of the Eager Dynamic Check

We can attempt to reason about the correctness of the eager dynamic check: all cases of eager scope extrusion are reported (by transitioning to **err**). This theorem is, unfortu-

¹Technically, the variable should be renamed. For clarity, I do not rename the variables in this example

nately, not true: we may cause eager scope extrusion, but never use the offending AST n in an unsafe way.

For example, recall that we can write non-terminating programs in λ_{op} (and thus $\lambda_{\langle\text{op}\rangle}$). Consider the $\lambda_{\langle\text{op}\rangle}$ program in Listing 15, where Ω is some non-terminating program that never refers to z . After translation, the program will reduce to a configuration that exhibits eager scope extrusion (**return** y). However, the program will immediately enter into a non-terminating loop, and scope extrusion will never be reported.

```
 $\$(\text{do } z \leftarrow (\text{handle} \langle \lambda x. \$(\text{extrude}(\langle \langle x \rangle \rangle); \langle \langle \text{return } x \rangle \rangle)) \rangle$   
   $\text{with } \{\text{return}(u) \mapsto \langle \langle 0 \rangle \rangle; \text{extrude}(y, k) \mapsto \text{return } y\}$   
   $\text{in } \Omega)$ 
```

$\lambda_{\langle\text{op}\rangle}$

Listing 15: The eager dynamic check is unsafe: it does not report all occurrences of eager scope extrusion. In the program, we cause eager scope extrusion (**return** y), where y refers to the unbound variable $\langle \langle x \rangle \rangle$ and then enter into a non-terminating loop which never refers to y .

Non-termination is not the only source of this behaviour. For example, we could throw away the offending AST (Listing 16), and therefore, never trigger the check:

```
 $\$(\text{handle} \langle \lambda x. \$(\text{extrude}(\langle \langle x \rangle \rangle); \langle \langle \text{return } x \rangle \rangle)) \rangle$   
   $\text{with } \{\text{return}(u) \mapsto \langle \langle 0 \rangle \rangle; \text{extrude}(y, k) \mapsto \text{do } w \leftarrow \text{return } y \text{ in } \langle \langle 0 \rangle \rangle\}$ 
```

$\lambda_{\langle\text{op}\rangle}$

Listing 16: The eager dynamic check additionally will not report eager scope extrusion in the case where the offending AST (**return** y , where y is bound to $\text{Var}(x_{\mathbb{N}})$) is thrown away.

Finally, in the most interesting example, we may *resume* the continuation. As shown in Listing 17, by resuming the continuation, we never use the offending AST $\text{Var}(x_{\mathbb{N}})$ in an *unsafe* way: w (bound to $\text{Var}(x_{\mathbb{N}})$) is only ever used in the context in which $\text{Var}(x_{\mathbb{N}})$ was declared safe:

```
 $\$(\text{handle} \langle \lambda x. \text{return } \$(\text{extrude}(\langle \langle x \rangle \rangle); \langle \langle \text{return } x \rangle \rangle)) \rangle$   
   $\text{with } \{\text{return}(u) \mapsto \text{return } u; \text{extrude}(y, k) \mapsto \text{do } w \leftarrow \text{return } y \text{ in } \text{continue } k \ w\}$ 
```

$\lambda_{\langle\text{op}\rangle}$

Listing 17: The eager dynamic check will additionally not report cases where the offending AST (**return** y , where y is bound to $\text{Var}(x_{\mathbb{N}})$) is used to build ASTs, but only in safe ways. In this case, the offending AST is bound to w , which is only ever used in a safe way: the continuation restores the context that permits $\text{Var}(x_{\mathbb{N}})$ to be used.

Expanding on Listing 17, **continue** returns the AST $\text{Var}(x_{\mathbb{N}})$ to the point where the **extrude** effect was performed, constructing the well-scoped AST: $\text{Lam}(\text{Var}(x_{\mathbb{N}}), \text{Var}(x_{\mathbb{N}}))$.

The permissiveness displayed in Listings 16 and 17 is arguably desirable: a *feature*, not a *bug* [14]. Permissiveness increases expressiveness, and is therefore desirable.

4.2.2 Expressiveness of the Eager Dynamic Check

Is the eager dynamic check maximally expressive? That is, does it only report errors that the lazy scope extrusion check would report? Listing 18 illustrates that the eager dynamic check is not maximally expressive, since we can use the offending AST in a way that *appears* to be unsafe, but is actually not. In Listing 18, the offending AST ($\text{Var}(x_{\mathbb{N}})$, bound to y) is used in a way that appears unsafe ($\langle\langle \$y + 0 \rangle\rangle$, bound to w), but the resulting AST w is *itself* not used in an unsafe way:

```
$(\text{handle } \langle\langle \lambda x. \$(\text{extrude}(\langle\langle x \rangle\rangle); \langle\langle \text{return } x \rangle\rangle) \rangle\rangle
  \text{with } \{ \text{return}(u) \mapsto \text{return } u; \text{extrude}(y, k) \mapsto \text{do } w \leftarrow \langle\langle \$y + 0 \rangle\rangle \text{ in continue } k w \})
```

$\lambda_{\langle\langle \text{op} \rangle\rangle}$

Listing 18: A program which will fail the eager scope extrusion check. The offending AST ($\text{Var}(x_{\mathbb{N}})$) is used to construct a larger AST in a way that appears to be unsafe $\langle\langle \$y + 0 \rangle\rangle$, but is actually not, since the unsafe AST is then only used in a safe way.

Expanding on Listing 18, the program fragment $\langle\langle \$y + 0 \rangle\rangle$ elaborates into

checkPlus($\text{Var}(x_{\mathbb{N}})$, $\text{Nat}(0)$)

and the **check** fails because $\text{Var}(x_{\mathbb{N}})$ is free and unbound. However, similarly to Listing 17, the AST $\text{Plus}(\text{Var}(x_{\mathbb{N}}), \text{Nat}(0))$ is returned to a context where $\text{Var}(x_{\mathbb{N}})$ is properly scoped.

The eager scope extrusion check is thus not as expressive as the lazy scope extrusion check: it will forbid some programs that the lazy scope extrusion check permits.

More concerningly, the eager dynamic check is unpredictable: it is difficult to explain, without appealing to the operational semantics, why the check disallows some programs, but allows others. Compare the program in Listing 17, which passes the check, and Listing 18, which *fails* the check. It is clear that the following inequation holds:

$$\langle\langle \$e \rangle\rangle \neq_s \langle\langle \$e + 0 \rangle\rangle$$

More generally, for program fragments P and P' (which are distinct from evaluation contexts):

$$P[e] =_s P'[e] \not\Rightarrow \langle\langle P[\langle\langle e \rangle\rangle] \rangle\rangle =_s \langle\langle P'[\langle\langle e \rangle\rangle] \rangle\rangle$$

One has to appeal to the operational behaviour of the eager dynamic check to explain why these equations do not hold. In my opinion, relying on the operational behaviour exposes too much of the internals of the check to the user.

One possible attempt to make the eager dynamic check better behaved is to change the **s**-mode elaboration of **return**, such that each **return** elaborates into a **check**

$$\llbracket \text{return } v \rrbracket_s = \text{check } \llbracket v \rrbracket_s$$

While this restores certain equations (for example, Listing 16 will now report a scope extrusion error), it will break others. For example, it is no longer the case that

$$\text{do } _ \leftarrow \text{return } v \text{ in } e =_s e$$

since the **return** is elaborated into a **check**, which v may fail.

4.2.3 Efficiency of the Eager Dynamic Check

Additionally, the eager dynamic check is not, in the worst case, more efficient than the lazy dynamic check, since there exist pathological examples (Listing 14). Further, the overhead of the checks cannot be bound: multi-shot continuations may be replayed an unbounded number of times. Multi-shot continuations may capture checks. Thus, the overhead of checking is unbounded. However, I conjecture that in the *common case*, the eager dynamic check detects scope extrusion sufficiently early as to outweigh the checking overhead. This argument is an empirical one, and would require further research, outside the scope of the dissertation, to support.

4.3 Best-Effort Dynamic Check

The lazy dynamic check is too inefficient (Section 2.3.1, Page 22), and the eager dynamic check too unpredictable. Might it be possible to find a “goldilocks” solution: one that is more efficient than the lazy dynamic check, but more expressive and predictable than the eager dynamic check?

That is, such a check should allow the program in Listing 18, and should be permissive in a predictable way. For example, a check that detects, as eagerly as possible, programs that *must* cause lazy scope extrusion. I call this Best-Effort Scope Extrusion. To moderate expectations, I will not end up building such a check, but rather an approximation.

Definition 4.3.1 (Best-Effort Scope Extrusion) A $\lambda_{\text{AST}(\text{op})}$ configuration of the form

$$\langle t; E; U; M; I \rangle$$

exhibits best-effort scope extrusion if there exists some $i \in \mathbb{N}$ such that

$$\langle t; E; U; M; I \rangle \rightarrow^i \langle t'; E'; U'; M'; I' \rangle$$

and $\langle t'; E'; U'; M'; I' \rangle$ exhibits lazy scope extrusion.

I will build this check incrementally. Consider, again, the program in Listing 18. For clarity, I elaborate it into $\lambda_{\text{AST}(\text{op})}$ (with language extensions, like Plus, and some simplification). I have underlined the failing check.

```

handle
  do  $x \leftarrow \text{mkvar } \mathbb{N}$  in
    check(dlet( $x$ , do body  $\leftarrow$  (do  $a \leftarrow \text{binderToAST } x$  in extrude( $a$ ))
      in return Lam( $x$ , body)))
with
  {return( $u$ )  $\mapsto$  return  $u$ ;
  extrude( $y$ ,  $k$ )  $\mapsto$  do  $w \leftarrow$  checkPlus( $y$ , Nat(0)) in continue  $k$ ( $w$ )}
```

However, w is only used in a context where $\text{Var}(x_{\mathbb{N}})$ is declared safe

check fails, transitioning to err,
 since y is bound to $\text{Var}(x_{\mathbb{N}})$,
 which is not declared via dlet

The check fails because when an effect is performed, the variable $\text{Var}(x_{\mathbb{N}})$ is no longer declared safe. Since y is bound to $\text{Var}(x_{\mathbb{N}})$, we report an error when executing the han-

handler for **extrude**. The problem is that the continuation k can be used to restore a scope in which $\text{Var}(x_{\mathbb{N}})$ is safe. It is not clear, when the Plus AST is constructed and checked, that eager scope extrusion *must* lead to lazy scope extrusion.

To make the check more expressive, therefore, it may be useful to temporarily allow $\text{Var}(x_{\mathbb{N}})$ to extrude its scope, delaying error reporting until one knows for certain that one must have lazy scope extrusion.

To allow $\text{Var}(x_{\mathbb{N}})$ to temporarily extrude its scope, we may use the **check_M** primitive, which checks for scope extrusion, but turns a blind eye to some set of muted variables M . It suffices to add $\text{Var}(x_{\mathbb{N}})$ to M . Inspection of the (EFF-OP) rule (Figure 3.8, Page 38) shows that, when effects are performed, the variables which become unsafe (like $\text{Var}(x_{\mathbb{N}})$) are added to the set of muted variables

$$M \cup \pi_{\text{Var}}(E_2)$$

As an example, consider the reduction of:

```

handle
  dlet(Var( $x_{\mathbb{N}}$ ), do body  $\leftarrow$  extrude( $\text{Var}(x_{\mathbb{N}})$ ) in Lam( $\text{Var}(x_{\mathbb{N}})$ , body))
with
  {return( $u$ )  $\mapsto$  return  $u$ ;
   extrude( $y, k$ )  $\mapsto$  do  $w \leftarrow$  checkM Plus( $y, \text{Nat}(0)$ ) in continue  $k w$ }

```

which is a simplified version of Listing 18 (elaborated into $\lambda_{\text{AST}(\text{op})}$). Note additionally that the **check** has been replaced with a **check_M**. The first reduction step performs the operation, and the configuration steps to

```

do  $w \leftarrow$  checkM Plus( $\text{Var}(x_{\mathbb{N}})$ ,  $\text{Nat}(0)$ ) in continue  $k w$ 

```

Since the **dlet**(Var($x_{\mathbb{N}}$), [−]) frame is no longer on the stack, **check** Plus($\text{Var}(x_{\mathbb{N}})$, $\text{Nat}(0)$) would throw an error. However, performing the **extrude** operation additionally *mutes* $\text{Var}(x_{\mathbb{N}})$, and hence **check_M** Plus($\text{Var}(x_{\mathbb{N}})$, $\text{Nat}(0)$) would not throw an error.

When should a variable like $\text{Var}(x_{\mathbb{N}})$ be *unmuted*? When we **do not** have any bound continuations k that could make the variable safe again. As a safe approximation for this, I track the point where we do not have *any* bound continuations k . I do so by keeping track of the maximal length I of the stack E that was never captured by the handling of an effect (Figure 4.1).

As an example, consider the following program, which builds the AST of the constant 1 function, $\lambda z. (\lambda x. x + 0)(1)$:

```

checkM(dlet(Var( $z_{\mathbb{N}}$ ), do  $b \leftarrow$ 
  (do  $f \leftarrow$  handle
    dlet(Var( $x_{\mathbb{N}}$ ), do body  $\leftarrow$  extrude( $\text{Var}(x_{\mathbb{N}})$ ) in Lam( $\text{Var}(x_{\mathbb{N}})$ , body))
    with
    {return( $u$ )  $\mapsto$  return  $u$ ;
     extrude( $y, k$ )  $\mapsto$  do  $w \leftarrow$  checkM Plus( $y, \text{Nat}(0)$ ) in continue  $k w$ }
  in do  $a \leftarrow$  return  $\text{Nat}(1)$ 
  in checkM App( $f, a$ )
  in return Lam( $\text{Var}(z_{\mathbb{N}})$ ,  $b$ )))

```

Note that the body of f , coloured in grey, is the simplified version of Listing 18, elaborated into $\lambda_{\text{AST}(\text{op})}$. The **extrude** operation is caught by the handler (also in grey), and




Figure 4.1: An illustration of when variables are unmuted. The stack, E , grows rightwards. Effects are caught by handlers (**dark grey**). This captures a portion of the stack (**light grey**). We track the length of the stack, in **white**, that are never captured by an operation in this fashion. Frames in **white** will never be able to resume a continuation.

never captures the surrounding context (in **black**). The surrounding context, which must have no references to the captured continuation, k , is identified by I .

If the stack was never captured by the handling of an effect, for example, no operations were performed, then I is set to \top , $\forall n \in \mathbb{N}, \top \geq n$. Performing an effect can thus *decrease* I , but never increase it. This is the side condition on **EFF-OP**.

During reduction, when the length of the stack is less than, or equals to, I , there must not be any remaining references to any continuations k , and thus I may be reset to \top , and all muted variables may be unmuted. In the previous example, the program eventually reduces to

```

checkM(dlet(Var(zN), do b ←
  (do f ← [return Lam(Var(xN), Plus(Var(xN), Nat(0)))]
  in do a ← return Nat(1)
  in checkM App(f, a))
  in return Lam(Var(zN), b)))

```

Where the square brackets $[-]$ separate the evaluation context (outside) and the term (inside). At this point, since the length of the stack is less than or equals to I , it is safe to unmute all muted variables. When there are no muted variables, **check_M** and **check** have the same behaviour.

However, altering the semantics in such a manner means that any transition could potentially have a side effect: unmuting variables. To keep the semantics standard, and to more closely model the implementation of the check, I associate the act of unmuting with **dlet**. That is, the only configuration whose transition may unmute variables is

$$\langle \mathbf{dlet}(\underline{\text{Var}}(\alpha_T), \mathbf{return } n); E; U; M; I \rangle$$

In the previous example, the transition from **dlet**($\underline{\text{Var}}(z_N)$, **return** n) unmutes variables. Therefore, $\text{Var}(x_N)$ is still muted when we check the application constructor **check_M** $\text{App}(f, a)$, but unmuted when we check the constructor of the outer lambda. This is the side-effect/side-condition in the **SEC-DLT** rule (Figure 3.8).

The best-effort scope extrusion check is thus simple: change all **checks** to **check_Ms**. Since the context outside the top-level splice is *ambient* and *inert*, it can never be captured by a handler. Hence, the **check_M** at the top-level splice will always behave like a **check**. Additionally, since **check_Ms** are at least as permissive as **checks**, the best-effort dynamic check is at least as expressive as the eager dynamic check. I verify that the implementation of the best-effort scope extrusion check allows the program in Listing 18.

4.3.1 Correctness of the Best-Effort Dynamic Check

Correctness of the check is easy to show: if there is best-effort scope extrusion, then either one of the the **check**_{MS} fire, or none do. In the latter case, we degenerate to the lazy scope extrusion check, where the top-level-splice **check** fires.

Theorem 4.3.1 (Correctness of the Best-Effort Check) *If*

1. $\cdot \vdash_{\mathbf{c}} e : T^0 ! \emptyset; \emptyset,$
2. $\llbracket e \rrbracket_{\mathbf{c}}^{BE} = t$
3. $\langle \text{erase-checks}(t); [-]; \emptyset; \emptyset; \top \rangle$ *exhibits best-effort scope extrusion*

then there exists E, U, M, I such that

$$\langle t; [-]; \emptyset; \emptyset; \top \rangle \rightarrow^* \langle \mathbf{err}; E; U; M; I \rangle$$

Note that in the third condition we have to erase the checks in t , since adding a check can transform a program that will eventually exhibit lazy scope extrusion to one that will not (instead terminating early with an error).

4.3.2 Expressiveness of the Best-Effort Dynamic Check

We do not, however, have maximal expressiveness. The best-effort dynamic check occasionally misfires, reporting false positives, that the lazy dynamic check would not have reported. In particular, it does not allow the program in [Listing 19](#). The program in [Listing 19](#) attempts to build the (unsafe) AST $\lambda x. \mathbf{return} \ y$, where y has extruded its scope, but then throws it away, returning instead the AST of 1. Critically, the effect never throws past the constructor of the outer lambda, $\lambda x. [-]$. Hence, the program will eventually reduce to a configuration

$$\langle \mathbf{dlet}(\mathbf{Var}(x_{\mathbb{N}}), \mathbf{return} \ \mathbf{Lam}(\mathbf{Var}(x_{\mathbb{N}}), \mathbf{Var}(y_{\mathbb{N}}))); E :: \mathbf{check}[-]; U; \{\mathbf{Var}(y_{\mathbb{N}})\}; I \rangle$$

where $\text{len}(E :: \mathbf{check}[-]) < I$. Thus, the transition from this configuration will unmute $\mathbf{Var}(y_{\mathbb{N}})$, and the surrounding **check** will fail, as $\mathbf{Var}(y_{\mathbb{N}})$ is free, unbound, and unmuted.

```
$(\langle \lambda x. \$(\mathbf{handle} \langle \lambda y. \$(\mathbf{op}(y); \mathbf{return} \ y) \rangle)
  \mathbf{with} \{ \mathbf{return}(u) \mapsto \mathbf{return} \langle \langle 0 \rangle \rangle; \mathbf{op}(z, k) \mapsto \mathbf{return} \ z \} \rangle);
\langle \langle 1 \rangle \rangle
```

$\lambda_{\langle \mathbf{op} \rangle}$

Listing 19: The best-effort scope extrusion check reports false positives. This program attempts to build the (unsafe) AST $\lambda x. \mathbf{return} \ y$, where y has extruded its scope, but then throws it away, returning instead the AST of 1. Critically, the program will eventually unmute $\mathbf{Var}(y_{\mathbb{N}})$, and the surrounding **check** will fail, as $\mathbf{Var}(y_{\mathbb{N}})$ is free, unbound, and unmuted.

Instead, I prove what I term a “cause-for-concern” property. Effectively, it states that if the best-effort check fires, **and the offending AST n appears in the final result of**

compile-time computation, then there must be lazy scope extrusion. In other words: in cases like [Listing 19](#), the only valid thing to do with $\lambda x.\text{return } y$ is to throw it away.

Theorem 4.3.2 (Cause for Concern Property) *If*

1. $\cdot \vdash_{\mathbf{c}} e : T^0 ! \emptyset; \emptyset,$
2. $\llbracket e \rrbracket_{\mathbf{c}}^{BE} = t$
3. *There exists E, U, M, I such that $\langle t; [-]; \emptyset; \emptyset; \top \rangle \rightarrow^* \langle \text{check}_M n; E; U; M; I \rangle$ and $\langle \text{check}_M n; E; U; M; I \rangle \rightarrow \langle \text{err}; E; U; M; I \rangle$*

Then for all $i \in \mathbb{N}$, if

4. $\langle \text{return } n; \text{erase-checks}(E); U; M; I \rangle \rightarrow^i \langle \text{return } m; E'; U'; M'; I' \rangle$
5. *and n a subtree of m ,*

then $FV^0(m) \not\subseteq \pi_{\text{Var}}(E')$

The proof of [Theorem 4.3.2](#) is by contradiction. By assumption (3) there must be at least one variable, call it $\text{Var}(x_T)$, that is free in n and not properly declared in E . Assume for the sake of contradiction that $FV^0(m) \subseteq \pi_{\text{Var}}(E')$. Then, by assumption (5), the frame $\mathbf{dlet}(\text{Var}(x_T), [-])$ must be in E' . By assumptions (1) and (2), we only have to consider terms that are the target of elaboration. By definition of elaboration, the only way to push the frame $\mathbf{dlet}(\text{Var}(x_T), [-])$ on E' is by resuming a continuation containing $\mathbf{dlet}(\text{Var}(x_T), [-])$. But then we must have had access to the resumed continuation in E . In turn, this implies $I < \text{len}(E)$, and thus $\text{Var}(x_T)$ would not have been unmuted. Consequently, $\text{check}_M n$ would not have failed because of $\text{Var}(x_T)$. This contradicts assumption (3).

Note that the eager dynamic check does not have the Cause for Concern property, with [Listing 18](#) being a counter-example. Hence, I conclude that the best-effort dynamic check is **more** expressive, and more **predictably** expressive, than the eager dynamic check.

4.3.3 Efficiency of the Best-Effort Dynamic Check

The best-effort dynamic check is certainly less efficient than the eager dynamic check. Like the eager dynamic check, I conjecture that the best-effort check is more efficient in the common case than the lazy check, but I leave substantiating the argument to further research.

4.4 Refined Environment Classifiers

Recall the objective of this chapter: to evaluate $\lambda_{\langle \text{op} \rangle}$ as a common language for encoding, and evaluating, different policies for moderating scope extrusion. I have encoded three dynamic checks into $\lambda_{\langle \text{op} \rangle}$, and considered their correctness, expressiveness, and efficiency. The evaluation was comparative: I was able to conclude that the best-effort

Types (Refined Environment Classifiers)

 $\lambda_{\langle\text{op}\rangle}$

Level -1 Values $T^{-1} ::= \dots \mid (\text{Code}(T^0 ! \xi)^{\gamma})^{-1}$

Figure 4.2: Extending $\lambda_{\langle\text{op}\rangle}$ with refined environment classifiers. The only change is that Code types are now annotated with an environment classifier γ , highlighted in **red**.

check was more expressive than the eager dynamic check, but less expressive than the lazy dynamic check. The comparative evaluation was facilitated by $\lambda_{\langle\text{op}\rangle}$.

However, dynamic checks are not the only way to prevent scope extrusion. I now turn my attention to **static** checks, where the type system is responsible for preventing scope extrusion. One such static check is the method of Refined Environment Classifiers (Section 2.3.1, Page 23). Isoda et al. [10] introduce a calculus with algebraic effects and code combinators (rather than quotes and splices), whose interaction is moderated via refined environment classifiers. I will demonstrate that refined environment classifiers may be encoded, and therefore evaluated, in $\lambda_{\langle\text{op}\rangle}$.

I first show that it is possible to augment the $\lambda_{\langle\text{op}\rangle}$ type system with a **simplified** version of Isoda et al.'s type system, as follows:

1. $\lambda_{\langle\text{op}\rangle}$ types are straightforwardly extended, by annotating all level -1 Code types with a classifier (Figure 4.2).
2. Similarly, level 0 types are associated with a classifier. However, following Isoda et al. [10], this association is indirect: for level 0 types, classifiers are associated with the judgement rather than annotated onto the type (Figure 4.3).
3. Correspondingly, I will augment the $\lambda_{\text{AST}(\text{op})}$ type system such that all AST types are annotated with a classifier.

For ease of reasoning, I will find it helpful to define the notion of an **extended** source type.

Definition 4.4.1 (Extended source type) *An **extended** source type is either:*

1. *A level -1 type, for example, $(\text{Code}(\mathbb{N}^0)^{\gamma})^{-1}$ or*
2. *A level 0 type annotated with a classifier, for example, $\mathbb{N}^0(\gamma)$.*
3. *A level 0 binder type, which is a level 0 value type (like $\mathbb{N}^0(\gamma)$) annotated with an underline to indicate that it will be elaborated into a binder type, $(\underline{\mathbb{N}^0(\gamma)})$*

Contexts Γ must contain the least classifier γ_{\perp} (so the empty context is no longer permitted), and the definition of well-typed expression is adapted to reflect this

Definition 4.4.2 (Well-typed Expression, Refined Environment Classifiers) e is a well-typed expression if $\gamma_{\perp} \vdash^{\gamma_{\perp}} e : T^0 ! \emptyset; \emptyset$

Most typing rules are straightforwardly adapted, and I have listed key typing rules in Figure 4.3.

The **c|q-LAMBDA** rule corresponds to C-Abs in the refined environment classifier literature. Following Isoda et al., I restrict handlers and continuations to Code types; However, I simplify types and typing rules by eliminating polymorphism. To allow for let-insertion [36], as in Listing 20, Isoda et al.’s typing rules for handlers and continuations are polymorphic over the classifier. In Listing 20, the **genlet** effect is used to generate more efficient code: $(\lambda z. z + z)e$, where e is only evaluated once, rather than $e + e$. The continuation k is resumed under a binder (which introduces a classifier). Thus, for the program to be well-typed, the continuation k should be polymorphic over classifiers. Similarly, the type system demands that handlers are polymorphic over classifiers as well. For example, a more faithful transcription of their handler type would have the form:

$$\forall \gamma. ((\text{Code}(T_1 ! \xi)^{\gamma})^{-1} ! \Delta \implies (\text{Code}(T_2 ! \xi')^{\gamma})^{-1} ! \Delta')^{-1}$$

without the polymorphism $(\forall \gamma)$, it will be difficult to prove that programs like the one in Listing 20 will produce well-scoped ASTs.

```
$ (handle do  $x \leftarrow \text{genlet}(\langle\langle e \rangle\rangle)$  in  $\langle\langle x + x \rangle\rangle$ 
  with {return( $u$ )  $\mapsto u$ ;  $\text{genlet}(y, k) \mapsto \langle\langle (\lambda z. \$(\text{continue } k \langle\langle z \rangle\rangle)) y \rangle\rangle$ })
```

$\lambda_{\langle\langle \text{op} \rangle\rangle}$

Listing 20: An example of let-insertion. For the program to be well-typed, the continuation k should be polymorphic over classifiers. To prove that continuations are always polymorphic in this manner, the type system demands that handlers are polymorphic over classifiers as well.

Reasoning about the correctness of polymorphic typing rules, however, is complex. It is even more complex in $\lambda_{\langle\langle \text{op} \rangle\rangle}$. Recall that $\lambda_{\langle\langle \text{op} \rangle\rangle}$ does not have an operational semantics, but rather must first be elaborated into $\lambda_{\text{AST}(\text{op})}$ terms. Thus, it is difficult to reason directly about the $\lambda_{\langle\langle \text{op} \rangle\rangle}$ type system via progress and preservation. As I will later demonstrate, one has to appeal to alternative techniques, like Tait-style logical relations [30], increasing the complexity of reasoning.

Given that the focus of this evaluation is on $\lambda_{\langle\langle \text{op} \rangle\rangle}$, rather than refined environment classifiers, I choose to simplify the type system, and minimise the complexity of reasoning.

It is unclear whether Isoda et al.’s system allows for non-termination. Once again, to simplify reasoning, I will force effect signatures to be well-founded and, in particular, not recursive. As proven by Kammar et al. [12], all well-typed programs must thus terminate.

Following Isoda et al. [10], I prove a weakening lemma. As types are stratified into two levels, and into value, computation, and handler types, we actually have 6 sub-lemmas. I list two as examples.

Lemma 4.4.1 (Weakening for Level 0 Values) *If $\Gamma \vdash_{\mathbf{c}|\mathbf{q}}^{\gamma} e : T^0$ then*

1. $\Gamma, (x : T_1^0)^{\gamma'} \vdash_{\mathbf{c}|\mathbf{q}}^{\gamma} e : T^0$ for some $\gamma' \in \Gamma$
2. $\Gamma, (x : T_1^{-1}) \vdash_{\mathbf{c}|\mathbf{q}}^{\gamma} e : T^0$
3. $\Gamma, \gamma' \vdash_{\mathbf{c}|\mathbf{q}}^{\gamma} e : T^0$, for some $\gamma' \notin \Gamma$
4. $\Gamma, \gamma' \sqsubseteq \gamma'' \vdash_{\mathbf{c}|\mathbf{q}}^{\gamma} e : T^0$, for some $\gamma', \gamma'' \in \Gamma$

Lemma 4.4.2 (Weakening for Level -1 Values) *If $\Gamma \vdash_{\mathbf{s}} e : T^{-1}$ then*

1. $\Gamma, (x : T_1^0)^{\gamma'} \vdash_{\mathbf{s}} e : T^{-1}$ for some $\gamma' \in \Gamma$
2. $\Gamma, (x : T_1^{-1}) \vdash_{\mathbf{s}} e : T^{-1}$
3. $\Gamma, \gamma' \vdash_{\mathbf{s}} e : T^{-1}$, for some $\gamma' \notin \Gamma$
4. $\Gamma, \gamma' \sqsubseteq \gamma'' \vdash_{\mathbf{s}} e : T^{-1}$, for some $\gamma', \gamma'' \in \Gamma$

Proof of these lemmas is by induction on the typing derivation.

Having encoded refined environment classifiers into the source type system, I now consider the elaboration into $\lambda_{\text{AST}(\text{op})}$ and the reduction of $\lambda_{\text{AST}(\text{op})}$ terms. I make the following changes:

1. Elaboration of types is defined on **extended** source types ([Definition 4.4.1](#)). This allows elaboration to carry classifiers. For example, rather than elaborating \mathbb{N}^0 , I elaborate $\mathbb{N}^0(\gamma)$ into $\text{AST}(\mathbb{N})^{\gamma}$.
2. Elaboration of contexts is extended to carry proof-theoretic terms like γ and $\gamma \sqsubseteq \gamma'$.
3. Elaboration of top-level splice $\llbracket \$e \rrbracket_{\mathbf{c}}$ is adapted to not insert **tls**.
4. For simplicity, since refined environment classifiers do not require any of the dynamic scope extrusion checks, I consider the subset of $\lambda_{\text{AST}(\text{op})}$ without **check** (and **check_M**), **dlet**, **tls**, and **err**. $\lambda_{\text{AST}(\text{op})}$ configurations can thus be shortened to $\langle t; E; U \rangle$, by dropping the machinery for muting and unmuting variables.

To the best of my knowledge, this is the first explicit presentation of refined environment classifiers in a calculus with quotation and splices rather than code combinators.

4.4.1 Correctness of Refined Environment Classifiers

I now reason about the correctness of refined environment classifiers. The correctness of refined environment classifiers was previously informally justified in [Section 2.3.1](#) ([Page 23](#)). In this section, I show that $\lambda_{\langle \text{op} \rangle}$ can be used to formally reason about correctness, by proving that every well-typed $\lambda_{\langle \text{op} \rangle}$ term returns a well-scoped AST on termination ([Theorem 4.4.1](#)).

Refined Environment Classifiers Typing Rules

 $\lambda_{\langle \text{op} \rangle}$

Selected Rules

$$\begin{array}{c}
\text{(c|q-VAR)} \\
\frac{(x : T^0)^\gamma \in \Gamma}{\Gamma \vdash_{\text{c|q}}^\gamma x : T^0 ! \Delta}
\end{array}
\quad
\begin{array}{c}
\text{(c|q-LAMBDA)} \\
\frac{\Gamma, \gamma', \gamma \sqsubseteq \gamma', (x : T_1^0)^{\gamma'} \vdash_{\text{c|q}}^{\gamma'} e : T_2^0 ! \Delta; \xi \quad \gamma \in \Gamma \quad \gamma' \notin \Gamma}{\Gamma \vdash_{\text{c|q}}^\gamma \lambda x. e : (T_1^0 \xrightarrow{\xi} T_2^0)^0 ! \Delta}
\end{array}$$

$$\begin{array}{c}
\text{(s-OP)} \\
\frac{\Gamma \vdash_{\text{s}} v : T_1^{-1} \quad \text{op} : T_1^{-1} \rightarrow (\text{Code}(T_2 ! \xi)^\gamma)^{-1} \in \Sigma \quad \text{op} \in \Delta}{\Gamma \vdash_{\text{s}} \text{op}(v) : (\text{Code}(T_2 ! \xi)^\gamma)^{-1} ! \Delta}
\end{array}$$

$$\begin{array}{c}
\text{(s-CONTINUE)} \\
\frac{\Gamma \vdash_{\text{s}} v_1 : ((\text{Code}(T_1 ! \xi)^\gamma)^{-1} \xrightarrow{\Delta} (\text{Code}(T_2 ! \xi')^{\gamma'})^{-1})^{-1} \quad \Gamma \vdash_{\text{s}} v_2 : (\text{Code}(T_1 ! \xi)^\gamma)^{-1}}{\Gamma \vdash_{\text{s}} \text{continue } v_1 v_2 : (\text{Code}(T_2 ! \xi')^{\gamma'})^{-1} ! \Delta}
\end{array}$$

$$\begin{array}{c}
\text{(s-HANDLE)} \\
\frac{\Gamma \vdash_{\text{s}} e : (\text{Code}(T_1 ! \xi)^\gamma)^{-1} ! \Delta \quad \Gamma \vdash_{\text{s}} h : ((\text{Code}(T_1 ! \xi)^\gamma)^{-1} ! \Delta \implies (\text{Code}(T_2 ! \xi')^\gamma)^{-1} ! \Delta')^{-1} \quad \forall \text{op} \in \Delta \setminus \Delta'. \text{op} \in \text{dom}(h)}{\Gamma \vdash_{\text{s}} \text{handle } e \text{ with } \{h\} : (\text{Code}(T_2 ! \xi')^\gamma)^{-1} ! \Delta'}
\end{array}$$

$$\begin{array}{c}
\text{(c|q-SPLICE)} \\
\frac{\Gamma \vdash_{\text{s}} e : (\text{Code}(T^0 ! \xi)^\gamma)^{-1} ! \Delta}{\Gamma \vdash_{\text{c|q}}^\gamma \$e : T^0 ! \Delta; \xi}
\end{array}
\quad
\begin{array}{c}
\text{(s-QUOTE)} \\
\frac{\Gamma \vdash_{\text{c|q}}^\gamma e : T^0 ! \Delta; \xi}{\Gamma \vdash_{\text{s}} \langle\langle e \rangle\rangle : (\text{Code}(T^0 ! \xi)^\gamma)^{-1} ! \Delta}
\end{array}$$

$$\begin{array}{c}
\text{(c|q-SUB)} \\
\frac{\Gamma \vdash_{\text{c|q}}^\gamma \$e : T^0 ! \Delta; \xi \quad \Gamma \models \gamma \sqsubseteq \gamma'}{\Gamma \vdash_{\text{c|q}}^{\gamma'} e : T^0 ! \Delta; \xi}
\end{array}
\quad
\begin{array}{c}
\text{(s-SUB)} \\
\frac{\Gamma \vdash_{\text{s}} e : (\text{Code}(T^0 ! \xi)^\gamma)^{-1} ! \Delta \quad \Gamma \models \gamma \sqsubseteq \gamma'}{\Gamma \vdash_{\text{s}} e : (\text{Code}(T^0 ! \xi)^{\gamma'})^{-1} ! \Delta}
\end{array}$$

Figure 4.3: Selected typing rules for refined environment classifiers. The **c|q-LAMBDA** rule corresponds to C-Abs in the refined environment classifier literature. Following Isoda et al., I restrict handlers and continuations to Code types; However, I simplify types and typing rules by eliminating polymorphism.

Theorem 4.4.1 (Correctness of Refined Environment Classifiers) *If*

1. $\cdot \vdash_{\mathbf{c}}^{\gamma_{\perp}} e : T^0 ! \emptyset; \emptyset$, and
2. $\llbracket e \rrbracket_{\mathbf{c}} = t$

then for some U , $\langle t; [-]; \emptyset \rangle \rightarrow^ \langle \mathbf{return} \ n; [-]; U \rangle$, and $FV^0(n) = \emptyset$*

The proof of [Theorem 4.4.1](#) is via a Tait-style logical relation. A logical relation is useful to show that typing guarantees are maintained by elaboration [3].

The definition of the logical relation, which I call *Scoped*, is presented in [Figure 4.4](#). *Scoped* is defined on core language $(\lambda_{\text{AST}(\text{op})})$ terms, and is indexed by:

1. A context of proof-theoretic terms Θ . Given a context of proof theoretic terms and variables Γ , one can project out only the proof theoretic terms $\pi_{\gamma}(\Gamma)$. For example, given

$$\Gamma = \gamma_{\perp}, \gamma_1, \gamma_{\perp} \sqsubseteq \gamma_1, (x : \mathbb{N}^0)^{\gamma_1}, \gamma_2, \gamma_1 \sqsubseteq \gamma_2, y : (\text{Code}(\mathbb{N}^0 ! \emptyset)^{\gamma_2})^{-1}$$

the proof theoretic part of the context is

$$\pi_{\gamma}(\Gamma) = \gamma_{\perp}, \gamma_1, \gamma_{\perp} \sqsubseteq \gamma_1, \gamma_2, \gamma_1 \sqsubseteq \gamma_2$$

which serves as our Θ .

2. An **extended** source-level type ([Definition 4.4.1](#)).

Most definitions are standard. I highlight two important definitions: the logical relation on the $T^0(\gamma)$ value type, and the logical relation on terms.

For a normal form n to be in the relation of the $T^0(\gamma)$ type, n must be an AST of the right type **and** the free variables of n need to be permitted by γ (permissibility was previously defined in [Section 2.3.1, Page 23](#)). Recall that the definition of permissibility assumes some known partial order on classifiers, e.g. $\gamma' \sqsubseteq \gamma$. The partial order is carried by the index Θ .

The logical relation on terms is defined as a least fixed point, following the definitions by Plotkin and Xie [22] and Kuchta [18], where the well-foundedness of the definition is additionally justified. I rely on the following induction principle, which I call *Scoped-Induction*²

Definition 4.4.3 (Scoped-Induction) *For some property Φ on closed terms of type $\llbracket \tau ! \Delta \rrbracket$, if*

1. $\langle t; [-]; U \rangle \rightarrow^* \langle \mathbf{return} \ n; [-]; U' \rangle$ *implies* $\Phi(t)$
2. $\langle t; [-]; U \rangle \rightarrow^* \langle \mathbf{op}(n); E; U' \rangle \not\rightarrow$, with $\mathbf{op} : A^{-1} \rightarrow B^{-1}$, $n \in \text{Scoped}_{\Theta, A^{-1}}$, and for arbitrary $n' \in \text{Scoped}_{\Theta, B^{-1}}$, $\Phi(E[n'])$ *implies* $\Phi(t)$

Then for all $t \in \text{Scoped}_{\Theta, \tau ! \Delta}$, $\Phi(t)$

I additionally rely on a closure lemma [18]

²Plotkin and Xie call the induction principle *G*-induction

Lemma 4.4.3 (Closure under Anti-Reduction) *If $\langle t; E; U \rangle \rightarrow^* \langle t'; E'; U' \rangle$ and $E'[t'] \in \text{Scoped}_{\Theta, \tau! \Delta}$ then $E[t] \in \text{Scoped}_{\Theta, \tau! \Delta}$*

Finally, I define a notion of closed substitution $\rho \models \Gamma$. Care must be taken with substitution of level-0 variables since these should be in the logical relation for Binders rather than ASTs (note the second clause in [Definition 4.4.4](#)).

Definition 4.4.4 (Closed substitution) *Given a context Γ , and assuming $\Theta = \pi_\gamma(\Gamma)$, the set of closed substitutions $\rho \models \Gamma$ are defined inductively as follows:*

1. $() \models \gamma_\perp$
2. If $\rho \models \Gamma$ and $n \in \text{Scoped}_{\Theta, T^0(\gamma)}$ then $(\rho, n/x) \models \Gamma, (x : T^0)^\gamma$
3. If $\rho \models \Gamma$ and $n \in \text{Scoped}_{\Theta, T^{-1}}$ then $(\rho, n/x) \models \Gamma, (x : T^{-1})$
4. If $\rho \models \Gamma$ then $\rho \models \Gamma, \gamma$
5. If $\rho \models \Gamma$ then $\rho \models \Gamma, \gamma \sqsubseteq \gamma'$

It is now easy to show the fundamental lemma. Once again, because types are stratified, and typing judgements are indexed by a mode, the fundamental lemma decomposes into many sub-lemmas. One such is stated in [Lemma 4.4.4](#).

Lemma 4.4.4 (Fundamental Lemma $[\mathbf{c}, T^0! \Delta; \xi]$ of the Scoped Logical Relation) *If $\Gamma \vdash_{\mathbf{c}}^\gamma e : T^0! \Delta; \xi$ then for $\Theta = \pi_\gamma(\Gamma)$, and for all ρ such that $\rho \models \Gamma$,*

$$\llbracket e \rrbracket_{\mathbf{c}}(\rho) \in \text{Scoped}_{\Theta, T^0! \Delta; \xi(\gamma)}$$

Proof of [Lemma 4.4.4](#) is by induction on the source $\lambda_{\langle \text{op} \rangle}$ typing rules. I focus on an interesting case: the \mathbf{c} -LAMBDA (C-ABS) case, where (handwaving the effect on U for the sake of clarity) it suffices to show that for some arbitrary $\rho, \Gamma \models \rho$,

do $x \leftarrow \mathbf{mkvar} \text{ erase}(T_1^0(\gamma'))$ **in do** $\text{body} \leftarrow \llbracket e \rrbracket_{\mathbf{c}}(\rho)$ **in return** $\text{Lam}(x, \text{body})$

in $\text{Scoped}_{\Theta, (T_1^0 \xrightarrow{\xi} T_2^0)^0! \Delta(\gamma)}$. It is clear that we reduce to

do $\text{body} \leftarrow \llbracket e \rrbracket_{\mathbf{c}}(\rho, \text{Var}(\alpha_{T_1^{\gamma'}})/x)$ **in return** $\text{Lam}(\text{Var}(\alpha_{T_1^{\gamma'}}), \text{body})$

By anti-reduction ([Lemma 4.4.3](#)) it suffices to show that the term above is in the logical relation.

By weakening, and the induction hypothesis, $\llbracket e \rrbracket_{\mathbf{c}}(\rho, \text{Var}(\alpha_{T_1^{\gamma'}})/x) \in \text{Scoped}_{\Theta', T_2^0! \Delta; \xi(\gamma')}$, where $\Theta' = \Theta, \gamma', \gamma \sqsubseteq \gamma'$. I thus apply Scoped-Induction on $\llbracket e \rrbracket_{\mathbf{c}}(\rho, \text{Var}(\alpha_{T_1^{\gamma'}}))$. The only interesting case is the first:

1. If $\llbracket e \rrbracket_{\mathbf{c}}(\rho, \text{Var}(\alpha_{T_1^{\gamma'}})/x) \in \text{Scoped}_{\Theta', T_2^0! \Delta; \xi(\gamma')}$ reduces to some **return** n , then by the inductive hypothesis it is of AST type and all its free variables are permitted by γ' . The term **do** $\text{body} \leftarrow \text{return } n$ **in** $\text{Lam}(\text{Var}(\alpha_{T_1^{\gamma'}}), \text{body})$ reduces to $\text{Lam}(\text{Var}(\alpha_{T_1^{\gamma'}}), n)$,

where $\text{Var}(\alpha_{T_1^{\gamma'}})$ is bound. By the typing rules, of the free variables, α is the only variable tagged with classifier γ' . So we can conclude, under Θ , that the free variables of $\text{Lam}(\text{Var}(\alpha_{T_1^{\gamma'}}), n)$ are permitted by γ . The desired conclusion thus follows from anti-reduction.

2. If $\llbracket e \rrbracket_{\Theta}(\rho, \text{Var}(\alpha_{T_1^{\gamma'}})/x) \in \text{Scoped}_{\Theta', T_2^0! \Delta; \xi(\gamma')}$ reduces to $E[\text{op}(n)]$, then because the context **do** $\text{body} \leftarrow [-]$ **in return** $\text{Lam}(\text{Var}(\alpha_{T_1^{\gamma'}}), \text{body})$ introduces no handlers, the desired conclusion follows immediately from the inductive hypothesis.

[Theorem 4.4.1](#) is an immediate corollary of [Lemma 4.4.4](#) and [Definition 4.4.2](#).

I conjecture that this proof may be extended to support non-termination by incorporating the techniques of step-indexing and biorthogonality, as demonstrated by Bieracki et al. [4].

4.4.2 Expressiveness of Refined Environment Classifiers

The typing rules for refined environment classifiers forbid any program which attempts to extrude some variable x to a handler, *no matter how the handler chooses to use x* ([Listing 21](#)).

```
$(\text{handle } \langle \lambda x. \text{return } \$(\text{extrude}(\langle \langle x \rangle \rangle); \langle \langle \text{return } x \rangle \rangle) \rangle \rangle
  \text{with } \{ \text{return}(u) \mapsto \text{return } u; \text{extrude}(y, k) \mapsto \text{any arbitrary program} \})
```

$\lambda_{\langle \langle \text{op} \rangle \rangle}$

Listing 21: The typing rules for refined environment classifiers forbid any program which attempts to extrude some variable x to a handler. Even if x is unused, or is resumed safely, this program will not type check.

Thus, even if the handler throws x away, or resumes a continuation with x , the program in [Listing 21](#) will not type-check. Consequently, neither will [Listings 15](#) to [18](#). Thus, refined environment classifiers are less expressive even than the eager dynamic check.

The only variables that refined environment classifiers permit are those that will never cause scope extrusion, for example, the variable z in [Listing 22](#).

```
$\lambda z. (\text{handle } \langle \lambda x. \text{return } \$(\text{extrude}(\langle \langle z \rangle \rangle); \langle \langle \text{return } x \rangle \rangle) \rangle \rangle
  \text{with } \{ \text{return}(u) \mapsto \text{return } u; \text{extrude}(y, k) \mapsto \text{continue } k() \})
```

$\lambda_{\langle \langle \text{op} \rangle \rangle}$

Listing 22: Refined environment classifiers allow variables to be passed via effects, so long as the variable can never cause a scope extrusion error. In the program above, performing an effect with z will never cause a scope extrusion error.

4.5 Evaluation of $\lambda_{\langle \langle \text{op} \rangle \rangle}$

I thus conclude that the calculus $\lambda_{\langle \langle \text{op} \rangle \rangle}$ is an appropriate language in which to encode, and evaluate, scope extrusion checks. Formalising scope extrusion in $\lambda_{\langle \langle \text{op} \rangle \rangle}$ provided clarity, resulting in a new Best-Effort check, which I argue finds a sweet spot between

The $\text{Scoped}_{\Theta, T}$ Logical Relation

 $\lambda_{\langle \text{op} \rangle}$

Normal Forms

$$\begin{aligned}
n \in \text{Scoped}_{\Theta, \mathbb{N}^{-1}} &\triangleq n \in \mathbb{N} \\
n \in \text{Scoped}_{\Theta, T^0(\gamma)} &\triangleq \cdot \vdash \llbracket n \rrbracket \in \llbracket T^0(\gamma) \rrbracket \text{ and } \Theta \vdash \text{FV}^0(n) \subseteq \text{permitted}(\gamma) \\
&\quad (\text{symmetric for } (\text{Code}(T^0)^\gamma)^{-1}, T^0! \xi(\gamma), \text{ etc}) \\
n \in \text{Scoped}_{\Theta, T^0(\gamma)} &\triangleq \text{binderToAST } n \in \text{Scoped}_{\Theta, T^0! \Delta(\gamma)} \\
n \in \text{Scoped}_{\Theta, (T_1^{-1} \Delta \Rightarrow T_2^{-1})^{-1}} &\triangleq \forall n' \in \text{Scoped}_{\Theta, T_1^{-1}}, n' \in \text{Scoped}_{\Theta, T_2^{-1}! \Delta} \\
n \in \text{Scoped}_{\Theta, (T_1^{-1} \Delta \Rightarrow T_2^{-1})^{-1}} &\triangleq \forall n' \in \text{Scoped}_{\Theta, T_1^{-1}}, \text{continue } n' \in \text{Scoped}_{\Theta, T_2^{-1}! \Delta}
\end{aligned}$$

Handlers

$$\begin{aligned}
h \in \text{Scoped}_{\Theta, (T_1^{-1}! \Delta \Rightarrow T_2^{-1}! \Delta')^{-1}} &\triangleq \text{if } h = \text{return}(x) \mapsto t_{\text{ret}} \\
&\quad \forall n' \in \text{Scoped}_{\Theta, T_1^{-1}}, t_{\text{ret}}[n'/x] \in \text{Scoped}_{\Theta, T_2^{-1}! \Delta'} \\
&\text{else } h = h'; \text{op}(x, k) \mapsto t_{\text{op}}, \text{op} : A^{-1} \rightarrow B^{-1} \\
&\quad h' \in \text{Scoped}_{\Theta, (T_1^{-1}! \Delta \Rightarrow T_2^{-1}! \Delta')^{-1}} \text{ and} \\
&\quad \forall n \in \text{Scoped}_{\Theta, A^{-1}}, n' \in \text{Scoped}_{\Theta, B^{-1} \Delta' \Rightarrow T_2^{-1}}, \\
&\quad t_{\text{op}}[n/x, n'/k] \in \text{Scoped}_{\Theta, T_2^{-1}! \Delta'}
\end{aligned}$$

Terms

In the following, let $\tau! \Delta$ be shorthand for any of $T^0! \Delta(\gamma)$, $T^0! \Delta; \xi(\gamma)$, $(T_1^0! \xi \Rightarrow T_2^0! \xi')^0! \Delta(\gamma)$, or $T^{-1}! \Delta$

$\text{Scoped}_{\Theta, \tau! \Delta} \triangleq$ The smallest property on terms t such that

1. For arbitrary U consistent with t , exists U' such that $\langle t; [-]; U \rangle \rightarrow^* \langle \text{return } n; [-]; U' \rangle$, and $n \in \text{Scoped}_{\Theta, \tau}$
2. For arbitrary U consistent with t , exists U' such that $\langle t; [-]; \emptyset \rangle \rightarrow^* \langle \text{op}(n); E; U' \rangle \not\rightarrow$, and
 - a) $\text{op} : A^{-1} \rightarrow B^{-1}$,
 - b) $n \in \text{Scoped}_{\Theta, A^{-1}}$, and
 - c) for all $n' \in \text{Scoped}_{\Theta, B^{-1}}$, $E[n'] \in \text{Scoped}_{\Theta, \tau! \Delta}$

Where, in this context, consistent with t means that for all $\text{Var}(\alpha_T)$ or $\text{Var}(\alpha_T) \in t$, $\alpha \in U$. Note further that the two possibilities are mutually exclusive.

Figure 4.4: The definition of the Scoped logical relation. Most definitions are standard. The logical relation on terms is defined as a least fixed point, following the definitions by Plotkin and Xie [22] and Kuchta [18], where the well-foundedness of the definition is additionally justified.

the eager and dynamic checks. Additionally, by unifying checks under a common language, I allow for comparative evaluation, developing a bank of $\lambda_{\langle\langle\text{op}\rangle\rangle}$ programs, which may serve as litmus tests of expressiveness.

The cost of encoding static and dynamic checks into the same language is that reasoning about the correctness of static checks becomes more complicated, since one can no longer prove correctness via progress and preservation. I argue that this is a necessary artefact of a language that attempts to unify static and dynamic checks, since dynamic checks should be inserted by elaboration. Further, I argue that it is a reasonable cost to pay for a comprehensive and comparative evaluation.

5 Conclusion

In this thesis, I have made significant progress on four fronts.

1. I build a novel two-stage calculus, $\lambda_{\langle\text{op}\rangle}$, that allows for effect handlers at both stages: compile-time and run-time. I prove standard metatheoretic results about $\lambda_{\langle\text{op}\rangle}$.
2. I formally describe and evaluate a range of scope extrusion solutions in $\lambda_{\langle\text{op}\rangle}$, including the existing MacoCaml check [14] and the refined environment classifiers approach [17]. Such comparison was facilitated by $\lambda_{\langle\text{op}\rangle}$, validating its design.
3. I formally describe and evaluate a novel Best-Effort check, which I prove is correct, and argue by comparison with other checks, occupies a goldilocks zone between expressiveness and efficiency.
4. I implement all three dynamic checks (lazy, eager, and best-effort) in MacoCaml.

5.1 Limitations and Future Work

Let-Insertion

This thesis was primarily concern with the *undesirable* interaction of metaprogramming and effects, rather than their *desirable* interaction: the previously described let-insertion [36].

I argue that $\lambda_{\langle\text{op}\rangle}$ is an ideal target in which to study let-insertion. Let-insertion has typically been studied in calculi where only pure programs can be generated [10]. Let-insertion is more interesting when one can generate effectful programs. Since the order of operation becomes more important, it can be more challenging to describe an “optimal” insertion point.

Additionally, I wish to extend my proof-of-correctness for refined environment classifiers to include polymorphism, and thus permit a certain interface for let-insertion.

Empirical Studies

I asserted that the dynamic checks differ in efficiency, but in the common case, rather than worst case. I would like to verify this with empirical studies. Kiselyov [14] additionally considers the *usability* of scope extrusion solutions, in terms of the informativeness of error messages. I did not consider this in my evaluation, but would like to perform user studies to more thoroughly evaluate usability.

Formalisation

The proofs in this thesis are pen-and-paper proofs. However, several are quite intricate (for example, the fundamental lemma of the Scoped logical relation). I would like to formalise the proofs in a theorem prover, like Agda.

Bibliography

- [1] D. Abrahams and A. Gurtovoy. *C++ Template Metaprogramming: Concepts, Tools, and Techniques from Boost and Beyond (C++ in Depth Series)*. Addison-Wesley Professional, 2004. ISBN 0321227255.
- [2] A. Bauer and M. Pretnar. An effect system for algebraic effects and handlers. *Logical Methods in Computer Science*, Volume 10, Issue 4, Dec. 2014. ISSN 1860-5974. doi: 10.2168/lmcs-10(4:9)2014. URL [http://dx.doi.org/10.2168/LMCS-10\(4:9\)2014](http://dx.doi.org/10.2168/LMCS-10(4:9)2014).
- [3] N. Benton and C.-K. Hur. Biorthogonality, step-indexing and compiler correctness. In *Proceedings of the 14th ACM SIGPLAN International Conference on Functional Programming, ICFP '09*, page 97–108, New York, NY, USA, 2009. Association for Computing Machinery. ISBN 9781605583327. doi: 10.1145/1596550.1596567. URL <https://doi.org/10.1145/1596550.1596567>.
- [4] D. Biernacki, M. Piróg, P. Polesiuk, and F. Sieczkowski. Handle with care: relational interpretation of algebraic effects and handlers. *Proc. ACM Program. Lang.*, 2(POPL), Dec. 2017. doi: 10.1145/3158096. URL <https://doi.org/10.1145/3158096>.
- [5] C. Calcagno, E. Moggi, and W. Taha. Closed types as a simple approach to safe imperative multi-stage programming. In U. Montanari, J. D. P. Rolim, and E. Welzl, editors, *Automata, Languages and Programming*, pages 25–36, Berlin, Heidelberg, 2000. Springer Berlin Heidelberg. ISBN 978-3-540-45022-1.
- [6] T.-J. Chiang, J. Yallop, L. White, and N. Xie. Staged compilation with module functors. *Proc. ACM Program. Lang.*, 8(ICFP), Aug. 2024. doi: 10.1145/3674649. URL <https://doi.org/10.1145/3674649>.
- [7] M. Felleisen and D. P. Friedman. Control operators, the secd-machine, and the λ -calculus. In M. Wirsing, editor, *Formal Description of Programming Concepts - III: Proceedings of the IFIP TC 2/WG 2.2 Working Conference on Formal Description of Programming Concepts - III, Ebberup, Denmark, 25-28 August 1986*, pages 193–222. North-Holland, 1987.
- [8] D. Hillerström and S. Lindley. Shallow effect handlers. In S. Ryu, editor, *Programming Languages and Systems*, pages 415–435, Cham, 2018. Springer International Publishing. ISBN 978-3-030-02768-1.
- [9] J. Inoue and W. Taha. Reasoning about multi-stage programs. In H. Seidl, editor, *Programming Languages and Systems*, pages 357–376, Berlin, Heidelberg, 2012. Springer Berlin Heidelberg. ISBN 978-3-642-28869-2.
- [10] K. Isoda, A. Yokoyama, and Y. Kameyama. Type-safe code generation with algebraic effects and handlers. In *Proceedings of the 23rd ACM SIGPLAN International*

- Conference on Generative Programming: Concepts and Experiences*, GPCE '24, page 53–65, New York, NY, USA, 2024. Association for Computing Machinery. ISBN 9798400712111. doi: 10.1145/3689484.3690731. URL <https://doi.org/10.1145/3689484.3690731>.
- [11] Jax-ML. Using grad on vmap on map on function containing sinc results in error. URL <https://github.com/jax-ml/jax/issues/10750>.
 - [12] O. Kammar, S. Lindley, and N. Oury. Handlers in action. In *Proceedings of the 18th ACM SIGPLAN International Conference on Functional Programming*, ICFP '13, page 145–158, New York, NY, USA, 2013. Association for Computing Machinery. ISBN 9781450323260. doi: 10.1145/2500365.2500590. URL <https://doi.org/10.1145/2500365.2500590>.
 - [13] O. Kiselyov. Delimited control in ocaml, abstractly and concretely. *Theoretical Computer Science*, 435:56–76, 2012. ISSN 0304-3975. doi: <https://doi.org/10.1016/j.tcs.2012.02.025>. URL <https://www.sciencedirect.com/science/article/pii/S0304397512001661>. Functional and Logic Programming.
 - [14] O. Kiselyov. The design and implementation of ber metaocaml. In M. Codish and E. Sumii, editors, *Functional and Logic Programming*, pages 86–102, Cham, 2014. Springer International Publishing. ISBN 978-3-319-07151-0.
 - [15] O. Kiselyov. Generating c: Heterogeneous metaprogramming system description. *Science of Computer Programming*, 231:103015, 2024. ISSN 0167-6423. doi: <https://doi.org/10.1016/j.scico.2023.103015>. URL <https://www.sciencedirect.com/science/article/pii/S0167642323000977>.
 - [16] O. Kiselyov. Metaocaml: Ten years later: System description. In *Functional and Logic Programming: 17th International Symposium, FLOPS 2024, Kumamoto, Japan, May 15–17, 2024, Proceedings*, page 219–236, Berlin, Heidelberg, 2024. Springer-Verlag. ISBN 978-981-97-2299-0. doi: 10.1007/978-981-97-2300-3_12. URL https://doi.org/10.1007/978-981-97-2300-3_12.
 - [17] O. Kiselyov, Y. Kameyama, and Y. Sudo. Refined environment classifiers. In A. Igarashi, editor, *Programming Languages and Systems*, pages 271–291, Cham, 2016. Springer International Publishing. ISBN 978-3-319-47958-3.
 - [18] W. Kuchta. A proof of normalization for effect handlers, Sept. 2023. URL <https://icfp23.sigplan.org/details/hope-2023/4/A-proof-of-normalization-for-effect-handlers>. Seattle, Washington, United States.
 - [19] J. L. Lawall and O. Danvy. Continuation-based partial evaluation. *SIGPLAN Lisp Pointers*, VII(3):227–238, July 1994. ISSN 1045-3563. doi: 10.1145/182590.182483. URL <https://doi.org/10.1145/182590.182483>.
 - [20] G. Mainland. Explicitly heterogeneous metaprogramming with metahaskell. *SIGPLAN Not.*, 47(9):311–322, Sept. 2012. ISSN 0362-1340. doi: 10.1145/2398856.2364572. URL <https://doi.org/10.1145/2398856.2364572>.

- [21] L. Phipps-Costin, A. Rossberg, A. Guha, D. Leijen, D. Hillerström, K. Sivaramakrishnan, M. Pretnar, and S. Lindley. Continuing webassembly with effect handlers. *Proc. ACM Program. Lang.*, 7(OOPSLA2), Oct. 2023. doi: 10.1145/3622814. URL <https://doi.org/10.1145/3622814>.
- [22] G. Plotkin and N. Xie. Handling the selection monad (full version), 2025. URL <https://arxiv.org/abs/2504.03890>.
- [23] M. Pretnar. An introduction to algebraic effects and handlers. invited tutorial paper. *Electronic Notes in Theoretical Computer Science*, 319:19–35, 2015. ISSN 1571-0661. doi: <https://doi.org/10.1016/j.entcs.2015.12.003>. URL <https://www.sciencedirect.com/science/article/pii/S1571066115000705>. The 31st Conference on the Mathematical Foundations of Programming Semantics (MFPS XXXI).
- [24] H. G. Rice. Classes of recursively enumerable sets and their decision problems. *Transactions of the American Mathematical Society*, 74(2):358–366, 1953. ISSN 00029947, 10886850. URL <http://www.jstor.org/stable/1990888>.
- [25] A. D. Robison. Impact of economics on compiler optimization. In *Proceedings of the 2001 Joint ACM-ISCOPE Conference on Java Grande, JGI '01*, page 1–10, New York, NY, USA, 2001. Association for Computing Machinery. ISBN 1581133596. doi: 10.1145/376656.376751. URL <https://doi.org/10.1145/376656.376751>.
- [26] G. Scherer. Deciding equivalence with sums and the empty type. In *Proceedings of the 44th ACM SIGPLAN Symposium on Principles of Programming Languages, POPL '17*, page 374–386, New York, NY, USA, 2017. Association for Computing Machinery. ISBN 9781450346603. doi: 10.1145/3009837.3009901. URL <https://doi.org/10.1145/3009837.3009901>.
- [27] M. Servetto and E. Zucca. A meta-circular language for active libraries. In *Proceedings of the ACM SIGPLAN 2013 Workshop on Partial Evaluation and Program Manipulation, PEPM '13*, page 117–126, New York, NY, USA, 2013. Association for Computing Machinery. ISBN 9781450318426. doi: 10.1145/2426890.2426913. URL <https://doi.org/10.1145/2426890.2426913>.
- [28] T. Sheard and S. P. Jones. Template meta-programming for haskell. In *Proceedings of the 2002 ACM SIGPLAN Workshop on Haskell, Haskell '02*, page 1–16, New York, NY, USA, 2002. Association for Computing Machinery. ISBN 1581136056. doi: 10.1145/581690.581691. URL <https://doi.org/10.1145/581690.581691>.
- [29] K. Sivaramakrishnan, S. Dolan, L. White, T. Kelly, S. Jaffer, and A. Madhavapeddy. Retrofitting effect handlers onto ocaml. In *Proceedings of the 42nd ACM SIGPLAN International Conference on Programming Language Design and Implementation, PLDI 2021*, page 206–221, New York, NY, USA, 2021. Association for Computing Machinery. ISBN 9781450383912. doi: 10.1145/3453483.3454039. URL <https://doi.org/10.1145/3453483.3454039>.
- [30] W. W. Tait. Intensional interpretations of functionals of finite type i. *The Journal of Symbolic Logic*, 32(2):198–212, 1967. ISSN 00224812. URL <http://www.jstor.org/stable/2271658>.

- [31] L. Tratt. Domain specific language implementation via compile-time meta-programming. *ACM Transactions on Programming Languages and Systems (TOPLAS)*, 30(6):1–40, 2008.
- [32] J. Vandebon, J. G. F. Coutinho, W. Luk, and E. Nurvitadhi. Enhancing high-level synthesis using a meta-programming approach. *IEEE Transactions on Computers*, 70(12):2043–2055, 2021. doi: 10.1109/TC.2021.3096429.
- [33] H. Wickham. *Advanced R*. Chapman and Hall/CRC, 2019.
- [34] N. Xie, Y. Cong, K. Ikemori, and D. Leijen. First-class names for effect handlers. *Proc. ACM Program. Lang.*, 6(OOPSLA2), Oct. 2022. doi: 10.1145/3563289. URL <https://doi.org/10.1145/3563289>.
- [35] N. Xie, L. White, O. Nicole, and J. Yallop. Macocaml: Staging composable and compilable macros. *Proc. ACM Program. Lang.*, 7(ICFP), Aug. 2023. doi: 10.1145/3607851. URL <https://doi.org/10.1145/3607851>.
- [36] J. Yallop and O. Kiselyov. Generating mutually recursive definitions. In *Proceedings of the 2019 ACM SIGPLAN Workshop on Partial Evaluation and Program Manipulation, PEPM 2019*, page 75–81, New York, NY, USA, 2019. Association for Computing Machinery. ISBN 9781450362269. doi: 10.1145/3294032.3294078. URL <https://doi.org/10.1145/3294032.3294078>.
- [37] J. Yallop, N. Xie, and N. Krishnaswami. flap: A deterministic parser with fused lexing. *Proc. ACM Program. Lang.*, 7(PLDI), June 2023. doi: 10.1145/3591269. URL <https://doi.org/10.1145/3591269>.